

GROUP - 11

BIKE RENTAL COUNT PREDICTION : STASTICAL TECHNIQUES

Advanced Statistical Learning - II

INSTRUCTOR : CHING CHI YANG

ABHISHEK BOGA - U00884042

VISHNU CHOUNDUR - U00869233

MANASWINI KOMMAREDDY - U00881454

JOHNPAUL CHIEMELIE ANAMEGE - U00870359

SAI NARENDRA REDDY LAKKIREDDY - U00892852



CONTENTS

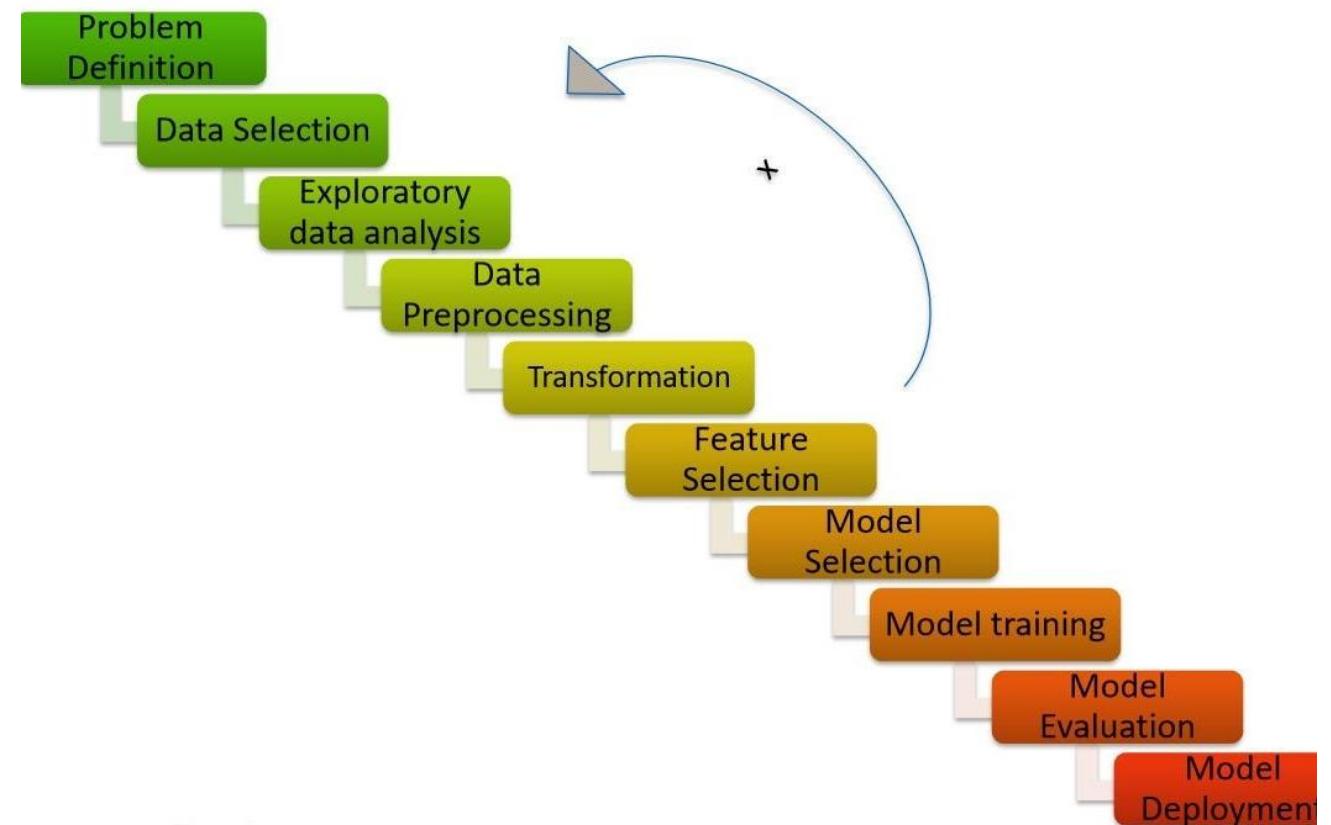
- Introduction
- Dataset
- Data Preprocessing
- Data Visualization
- Transformation
- Feature Selection
- Model building
 - ✓ Linear Regression
 - ✓ Linear Regression with Interaction Terms
 - ✓ Linear Regression with polynomial
 - ✓ Subset Selection – **Forward and Backward Selection**
 - ✓ Regularization Techniques or Shrinkage Methods – **Lasso and Ridge Regression**
 - ✓ Dimension Reduction Method – **PCR**
 - ✓ Random Forest Regression
 - ✓ Boosting Technique – **Gradient Boosting Machine**
 - ✓ Neural Network
- Results
- Conclusion
- Futurescope

INTRODUCTION

- At the heart of our project lies a singular goal: to craft a predictive model that not only reads the patterns of bike usage but also predicts them with precision. To this end, we are rigorously testing a spectrum of machine learning models, engaging in a comparative analysis to determine which algorithm reigns supreme in forecasting bike count based on variables like working days, seasons, and other vital parameters
- Beyond the aim of prediction accuracy, our mission is to meticulously evaluate and identify the most effective model from a cadre of contenders.
- Through this exploration, we aim to distill the essence of what makes a model not just functional, but exceptional, in predicting bike-sharing demand.

MACHINE LEARNING LIFE CYCLE

- **Problem Statement :** To predict the bike rental count based on hourly based on the environmental and seasonal settings.
- **Data Selection :**
 - Source : UC Irvine Machine Learning Repository,
 - Link :
<https://archive.ics.uci.edu/dataset/275/bike+sharing+dataset>
 - Here, two dataset we have : 1. hour.csv (17379*17) and day.csv (731*16).
 - We choose hour.csv dataset.
 - Dimension dataset : 17379 *17
 - 1 – Index, 2 – date, 3-10 – Categorical,
11-14 – Numerical and 15 – 17 – Dependent Variables



Data Set

- ✓ instant: record index
- ✓ dteday : date
- ✓ season : season (1:springer, 2:summer, 3:fall, 4:winter)
- ✓ yr : year (0: 2011, 1:2012)
- ✓ mnth : month (1 to 12)
- ✓ hr : hour (0 to 23)
- ✓ holiday : weather day is holiday or not
- ✓ weekday : day of the week
- ✓ workingday : if day is neither weekend nor holiday is 1, otherwise is 0.
- ✓ weathersit :
 - 1: Clear, Few clouds, Partly cloudy, Partly cloudy
 - 2: Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist
 - 3: Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rain + Scattered clouds
 - 4: Heavy Rain + Ice Pallets + Thunderstorm + Mist, Snow + Fog

instant	dteday	season	yr	mnth	hr	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual	registered	cnt	
1	2011-01-01		1	0	1	0	0	6	0	1	0.24	0.2879	0.81	0.0000	3	13	16
2	2011-01-01		1	0	1	1	0	6	0	1	0.22	0.2727	0.80	0.0000	8	32	40
3	2011-01-01		1	0	1	2	0	6	0	1	0.22	0.2727	0.80	0.0000	5	27	32
4	2011-01-01		1	0	1	3	0	6	0	1	0.24	0.2879	0.75	0.0000	3	10	13
5	2011-01-01		1	0	1	4	0	6	0	1	0.24	0.2879	0.75	0.0000	0	1	1
6	2011-01-01		1	0	1	5	0	6	0	2	0.24	0.2576	0.75	0.0896	0	1	1
7	2011-01-01		1	0	1	6	0	6	0	1	0.22	0.2727	0.80	0.0000	2	0	2

- ✓ temp : Normalized temperature in Celsius. The values are divided to 41 (max)
- ✓ atemp: Normalized feeling temperature in Celsius. The values are divided to 50 (max)
- ✓ hum: Normalized humidity. The values are divided to 100 (max)
- ✓ windspeed: Normalized wind speed. The values are divided to 67 (max)
- ✓ casual: count of casual users
- ✓ registered: count of registered users
- ✓ cnt: count of total rental bikes including both casual and registered

```

> summary(data)
      instant      dteday      season      yr      mnth      hr      holiday
Min. : 1 Length:17379   Min. :1.000   Min. :0.0000   Min. : 1.000   Min. : 0.00   Min. :0.00000
1st Qu.: 4346 Class :character 1st Qu.:2.000   1st Qu.:0.0000   1st Qu.: 4.000   1st Qu.: 6.00   1st Qu.:0.00000
Median : 8690 Mode  :character Median :3.000   Median :1.0000   Median : 7.000   Median :12.00   Median :0.00000
Mean   : 8690                   Mean :2.502   Mean :0.5026   Mean : 6.538   Mean :11.55   Mean :0.02877
3rd Qu.:13034                  3rd Qu.:3.000   3rd Qu.:1.0000   3rd Qu.:10.000  3rd Qu.:18.00  3rd Qu.:0.00000
Max.  :17379                  Max. :4.000   Max. :1.0000   Max. :12.000  Max. :23.00   Max. :1.00000
      weekday     workingday    weathersit      temp      atemp      hum      windspeed
Min. :0.000   Min. :0.0000   Min. :1.000   Min. :0.020   Min. :0.0000   Min. :0.0000   Min. :0.0000
1st Qu.:1.000  1st Qu.:0.0000 1st Qu.:1.000   1st Qu.:0.340   1st Qu.:0.3333   1st Qu.:0.4800   1st Qu.:0.1045
Median :3.000   Median :1.0000  Median :1.000   Median :0.500   Median :0.4848   Median :0.6300   Median :0.1940
Mean   :3.004   Mean :0.6827   Mean :1.425   Mean :0.497   Mean :0.4758   Mean :0.6272   Mean :0.1901
3rd Qu.:5.000   3rd Qu.:1.0000 3rd Qu.:2.000   3rd Qu.:0.660   3rd Qu.:0.6212   3rd Qu.:0.7800   3rd Qu.:0.2537
Max.  :6.000   Max. :1.0000   Max. :4.000   Max. :1.000   Max. :1.0000   Max. :1.0000   Max. :0.8507
      casual      registered      cnt
Min. : 0.00   Min. : 0.0   Min. : 1.0
1st Qu.: 4.00  1st Qu.: 34.0  1st Qu.: 40.0
Median :17.00  Median :115.0  Median :142.0
Mean   :35.68  Mean :153.8   Mean :189.5
3rd Qu.:48.00  3rd Qu.:220.0  3rd Qu.:281.0
Max.  :367.00  Max. :886.0   Max. :977.0

> dim(data) # 17379,17
[1] 17379 17

> str(data)
'data.frame': 17379 obs. of 17 variables:
 $ instant : int 1 2 3 4 5 6 7 8 9 10 ...
 $ dteday  : chr "2011-01-01" "2011-01-01" "2011-01-01" "2011-01-01" ...
 $ season  : int 1 1 1 1 1 1 1 1 1 ...
 $ yr      : int 0 0 0 0 0 0 0 0 0 ...
 $ mnth    : int 1 1 1 1 1 1 1 1 1 ...
 $ hr      : int 0 1 2 3 4 5 6 7 8 9 ...
 $ holiday : int 0 0 0 0 0 0 0 0 0 ...
 $ weekday : int 6 6 6 6 6 6 6 6 6 ...
 $ workingday: int 0 0 0 0 0 0 0 0 0 ...
 $ weathersit: int 1 1 1 1 1 2 1 1 1 ...
 $ temp    : num 0.24 0.22 0.22 0.24 0.24 0.24 0.22 0.2 0.24 0.32 ...
 $ atemp   : num 0.288 0.273 0.273 0.288 0.288 ...
 $ hum     : num 0.81 0.8 0.8 0.75 0.75 0.75 0.8 0.86 0.75 0.76 ...
 $ windspeed: num 0 0 0 0 0.0896 0 0 0 0 ...
 $ casual  : int 3 8 5 3 0 0 2 1 1 8 ...
 $ registered: int 13 32 27 10 1 1 0 2 7 6 ...
 $ cnt     : int 16 40 32 13 1 1 2 3 8 14 ...

> for (col_name in names(data)) {
+   if (length(unique(data[[col_name]])) < 25) {
+     cat(paste("Unique ", col_name, "'s count: ", length(unique(data[[col_name]])), "\n"))
+     cat(paste("Unique values: ", paste(unique(data[[col_name]])), collapse = ", "), "\n\n"))
+   }
+ }
Unique season 's count: 4
Unique values: 1, 2, 3, 4
Unique yr 's count: 2
Unique values: 0, 1
Unique mnth 's count: 12
Unique values: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12
Unique hr 's count: 24
Unique values: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23
Unique holiday 's count: 2
Unique values: 0, 1
Unique weekday 's count: 7
Unique values: 6, 0, 1, 2, 3, 4, 5
Unique workingday 's count: 2
Unique values: 0, 1
Unique weathersit 's count: 4
Unique values: 1, 2, 3, 4

```

Data Preprocessing

- Handling Missing Data
- Data Cleaning
- Handling Date and Time Data

```
> ## Duplicate Data
> # Find and display duplicate rows based on all columns
> duplicate_rows <- data[duplicated(data) | duplicated(data, fromLast = TRUE), ]
> # Display the duplicate rows
> print("Duplicate Rows in the Dataset:")
[1] "Duplicate Rows in the Dataset:"
> print(duplicate_rows)
   season yr mnth hr holiday weekday weathersit temp atemp hum windspeed cnt
7959     4  0    12  6      0       6        0  1  0.24  0.2576  0.65  0.1045  11
8127     4  0    12  6      0       6        0  1  0.24  0.2576  0.65  0.1045  11
13560    3  1     7  4      0       2        1  1  0.66  0.6061  0.83  0.0896   6
13728    3  1     7  4      0       2        1  1  0.66  0.6061  0.83  0.0896   6
... - . . . -
```

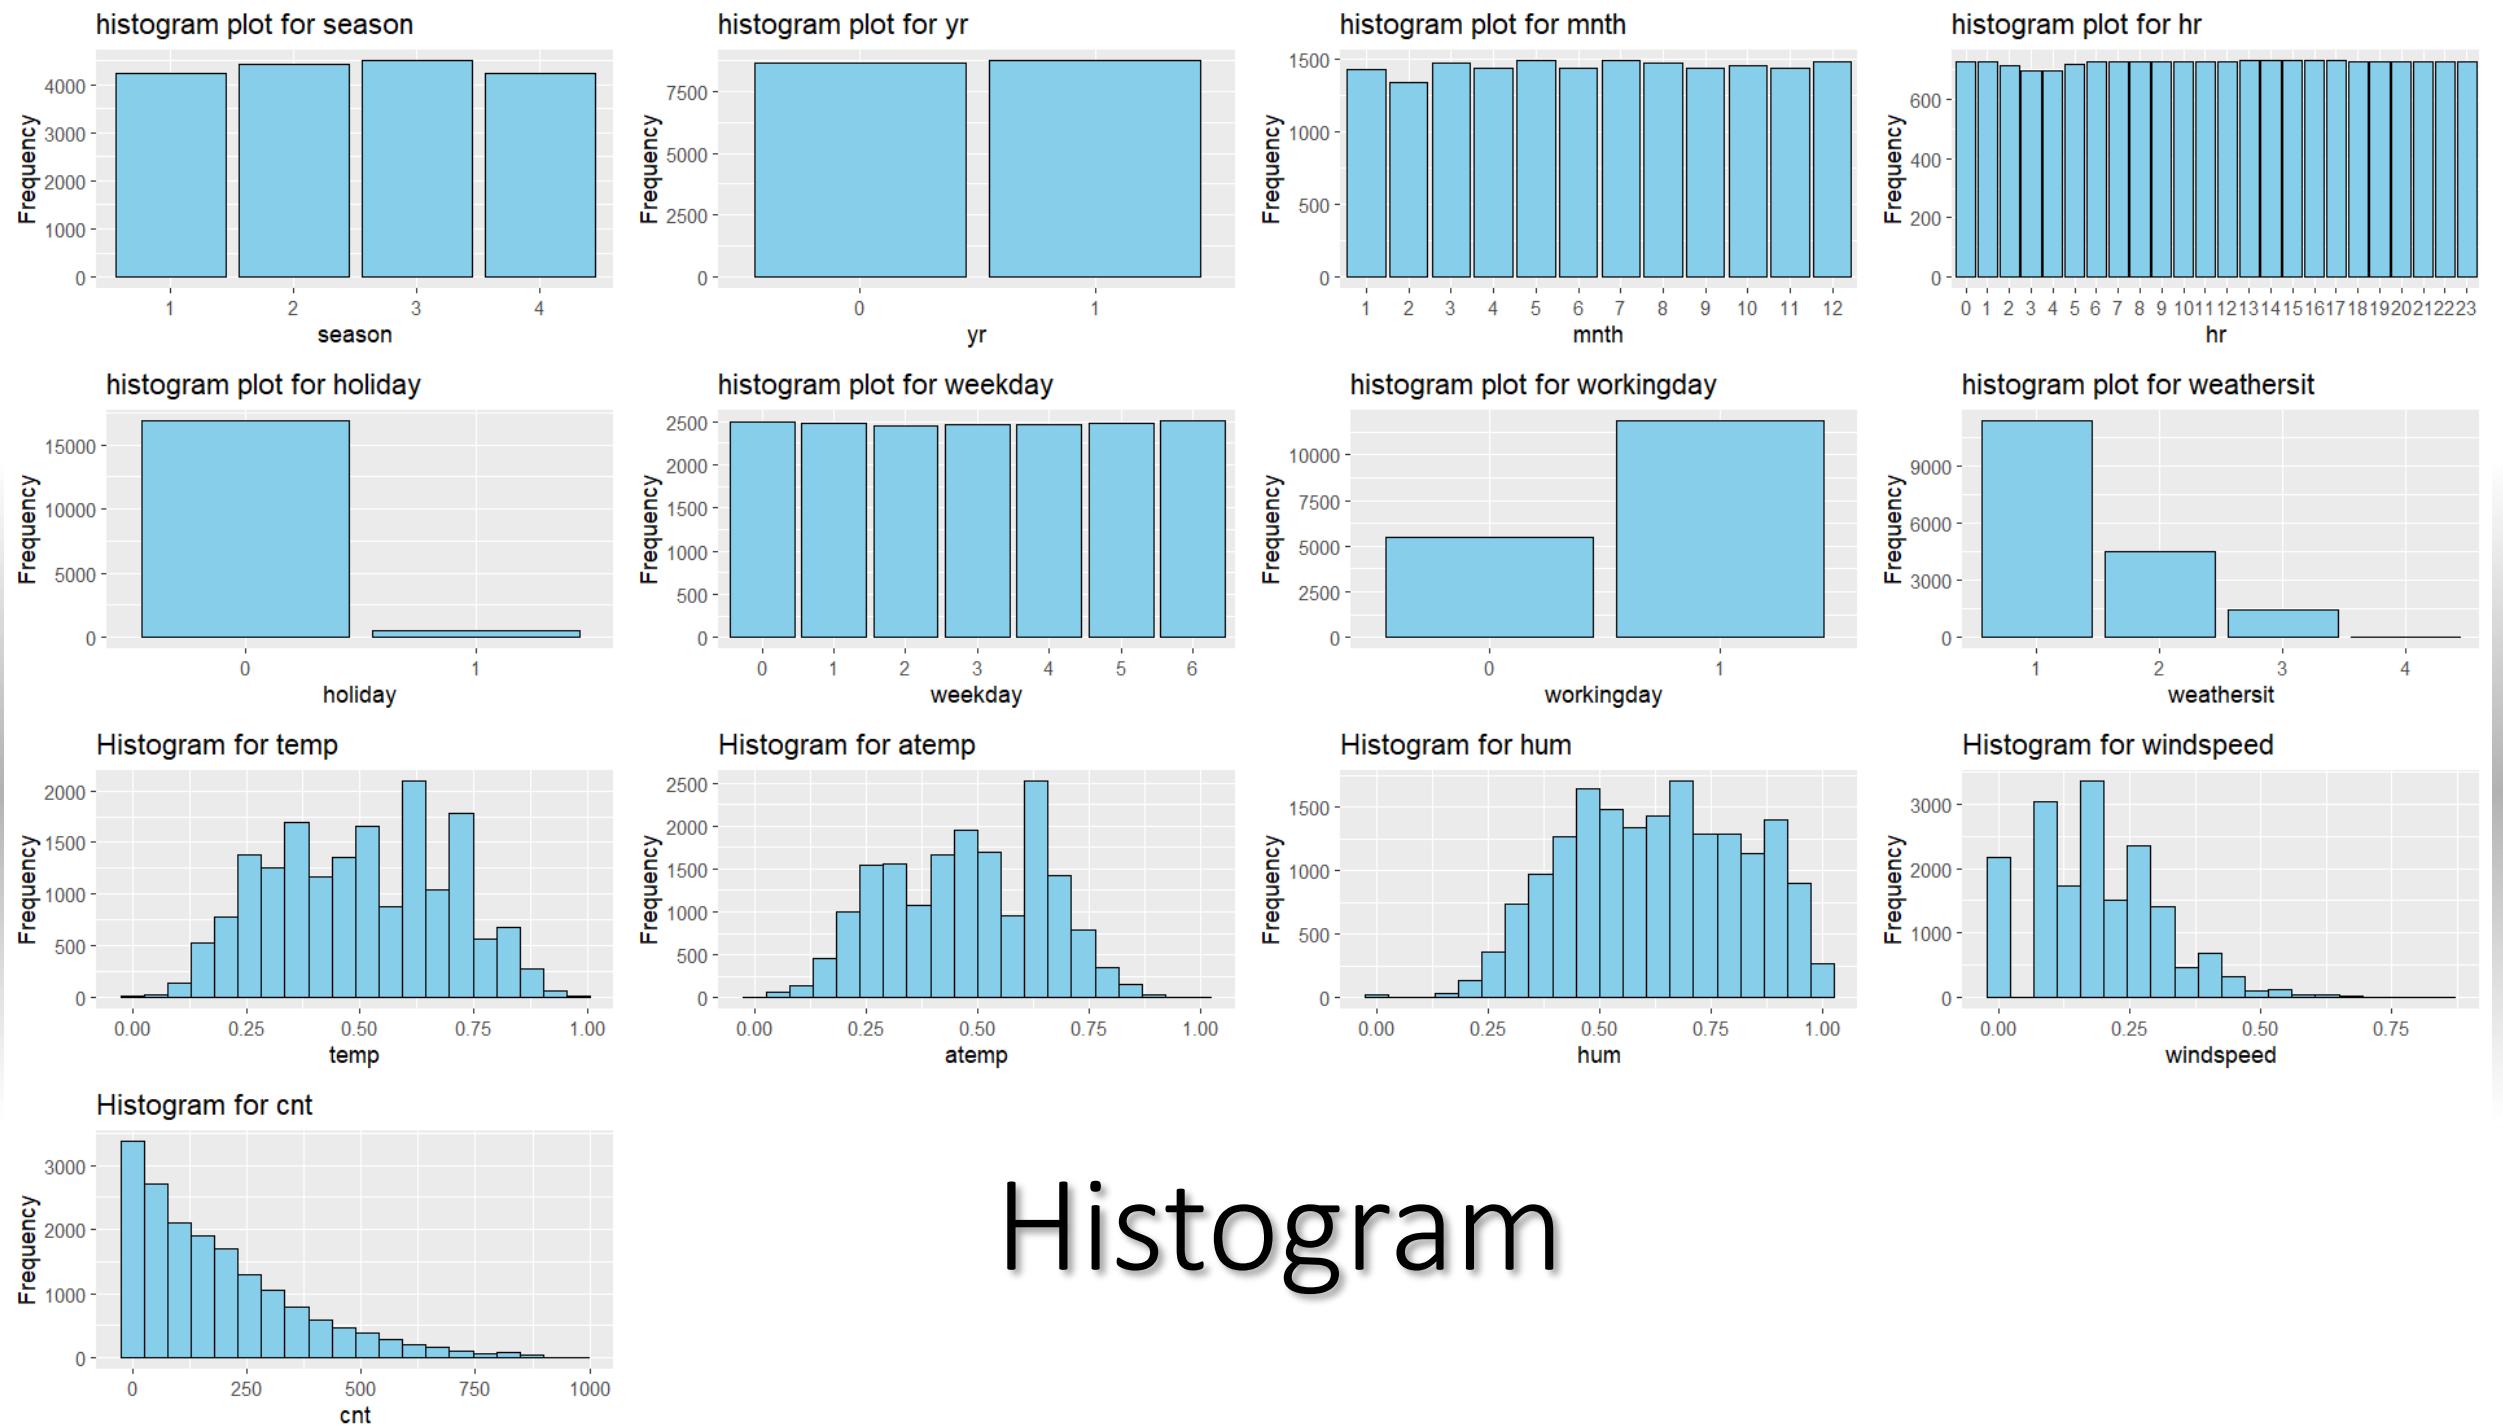
```
> #####
> ### Checking relationship between the dteday and yr
> data$Year <- format(as.Date(data$dteday), "%Y")
> data$Category <- ifelse(data$Year == "2011", 0, 1)
> identical_columns <- identical(data$yr, as.integer(data$Category))
> print(identical_columns)
[1] TRUE
> ### Checking relationship between the dteday and mnth
> data$month <- format(as.Date(data$dteday), "%m")
> identical_columns <- identical(data$mnth, as.integer(data$month))
> print(identical_columns)
[1] TRUE
> ### Checking relationship between the casual, registered and cnt
> data$cnt1 <- data$casual+data$registered
> identical_columns <- identical(data$cnt, data$casual+data$registered)
> print(identical_columns)
[1] TRUE
>
```

```
> print(missing_counts)
   season          yr      mnth      hr      holiday      weekday      workingday      weathersit      temp      atemp
           0           0       0       0           0           0           0           0           0           0
      hum      windspeed      cnt
           0           0       0
```

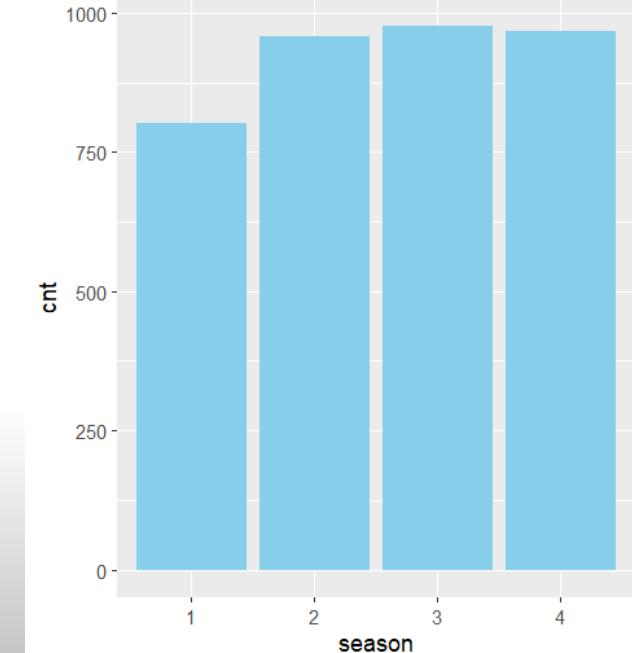


Data Visualization And Analysis

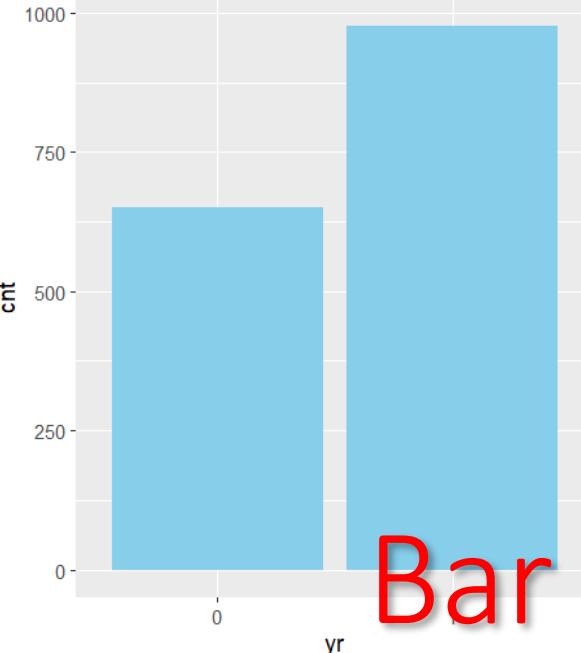
- Histogram
- Scatter Plot
- Box Plot
- Correlation



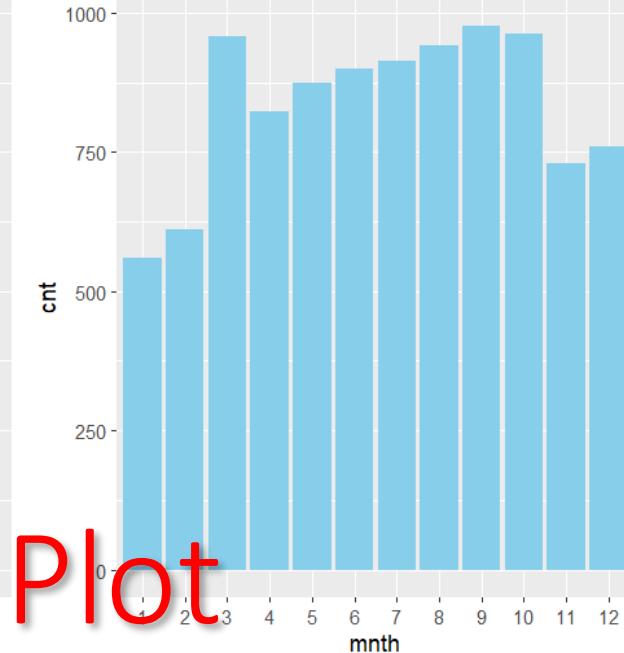
Bar Plot for season



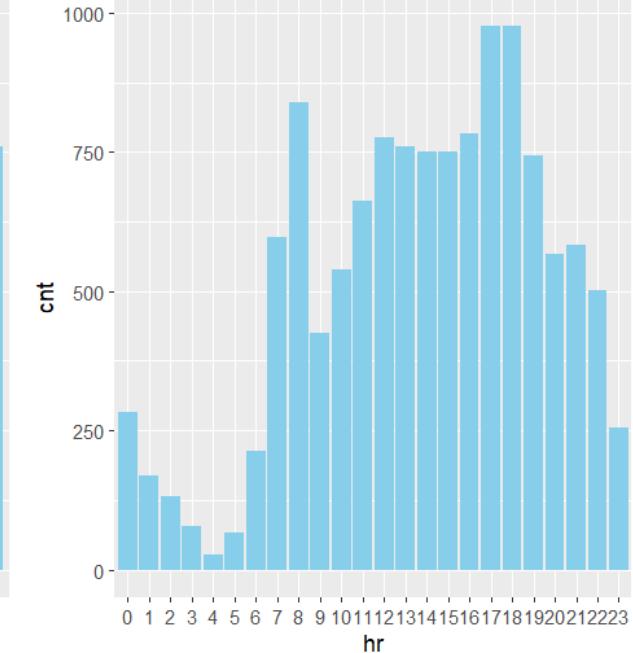
Bar Plot for yr



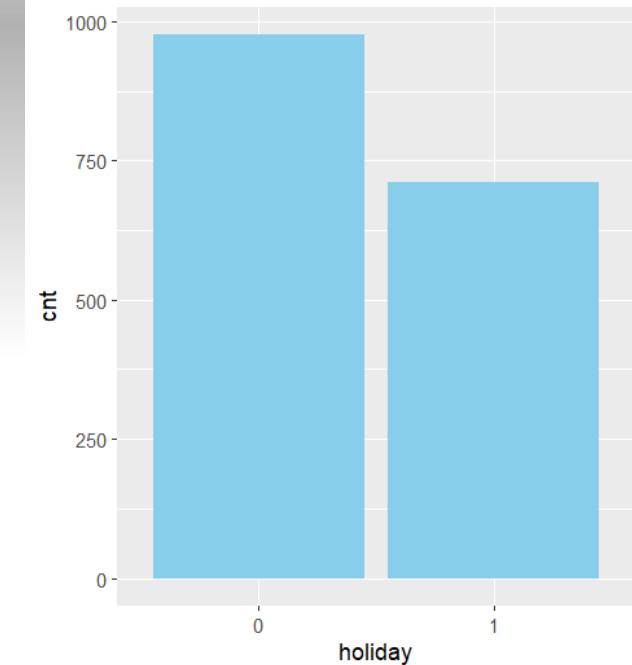
Bar Plot for mnth



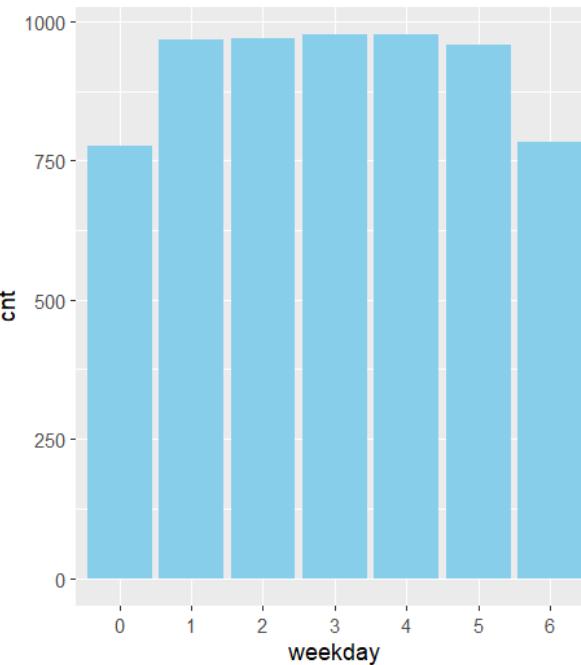
Bar Plot for hr



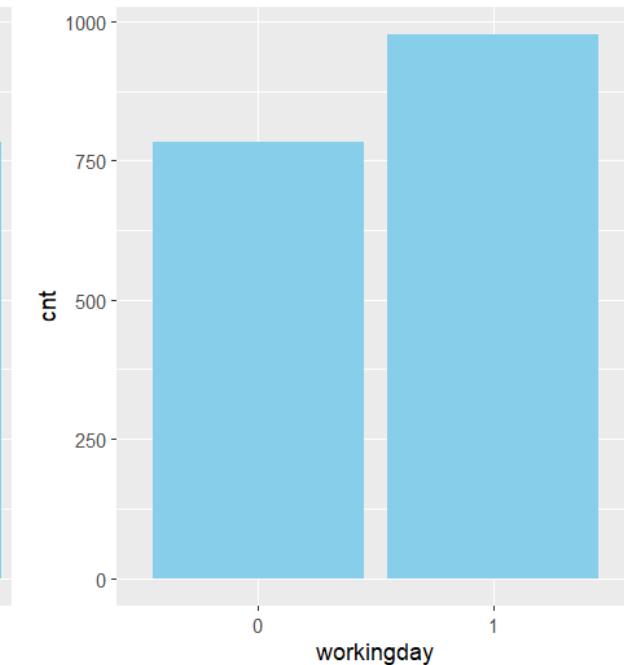
Bar Plot for holiday



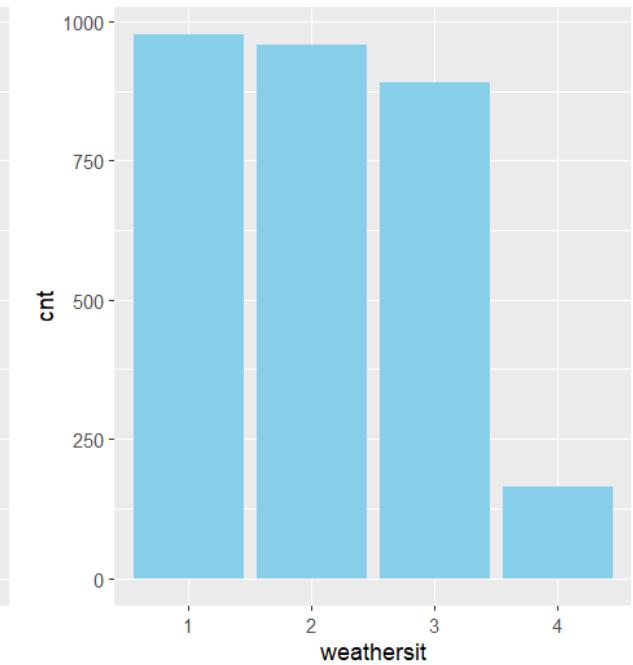
Bar Plot for weekday



Bar Plot for workingday

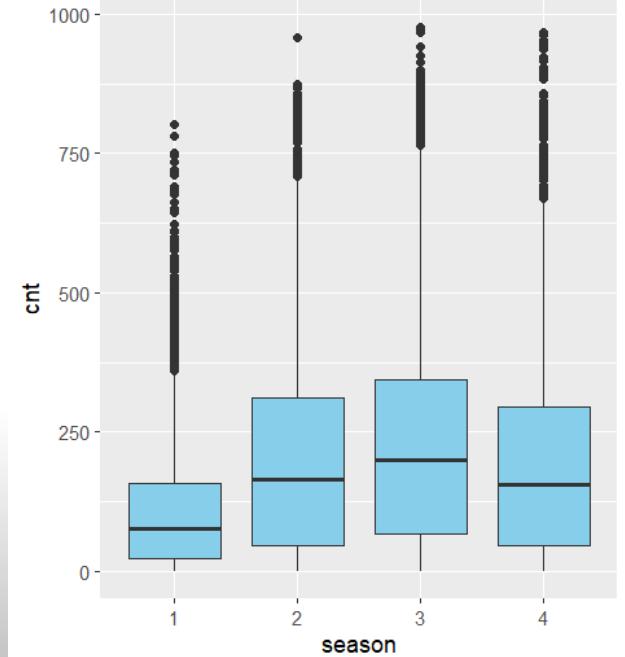


Bar Plot for weathersit

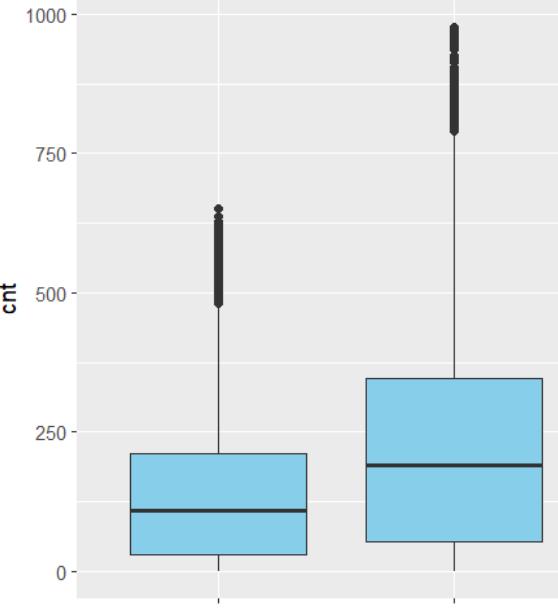


Bar Plot

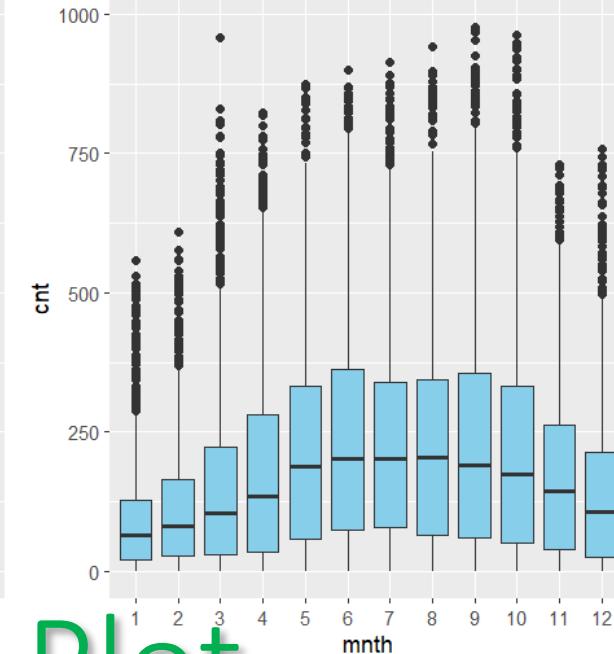
Box Plot for season



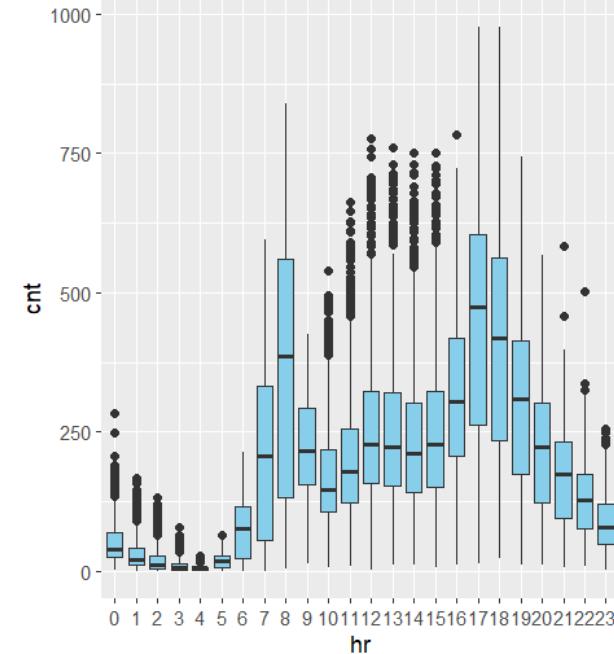
Box Plot for yr



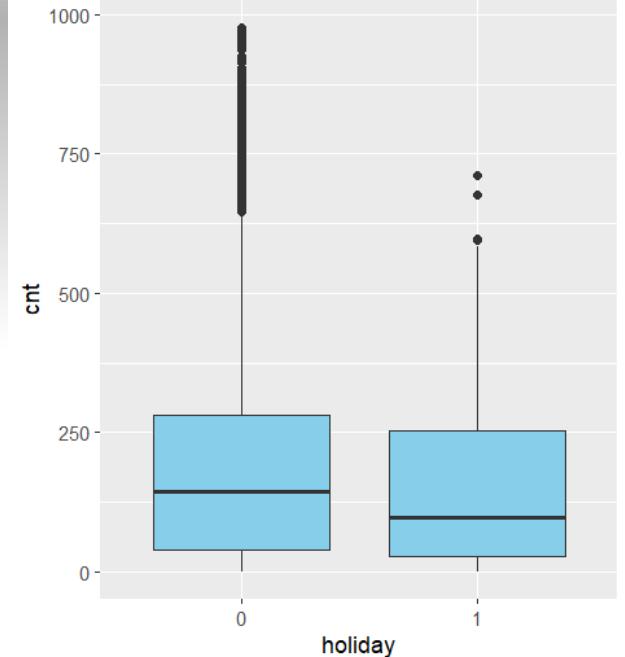
Box Plot for mnth



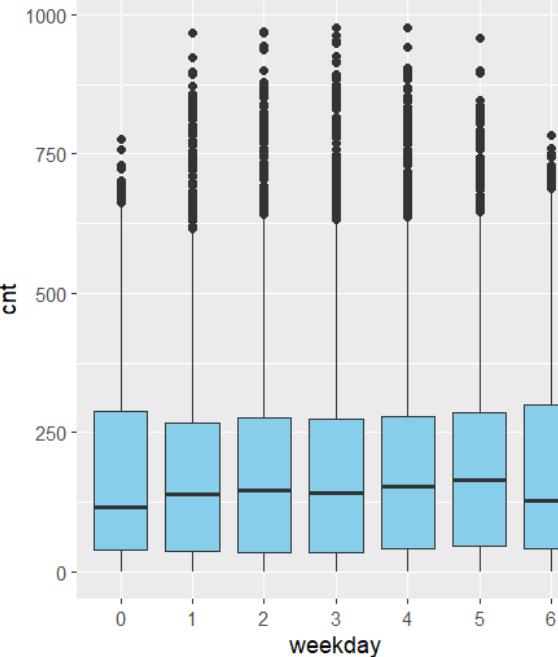
Box Plot for hr



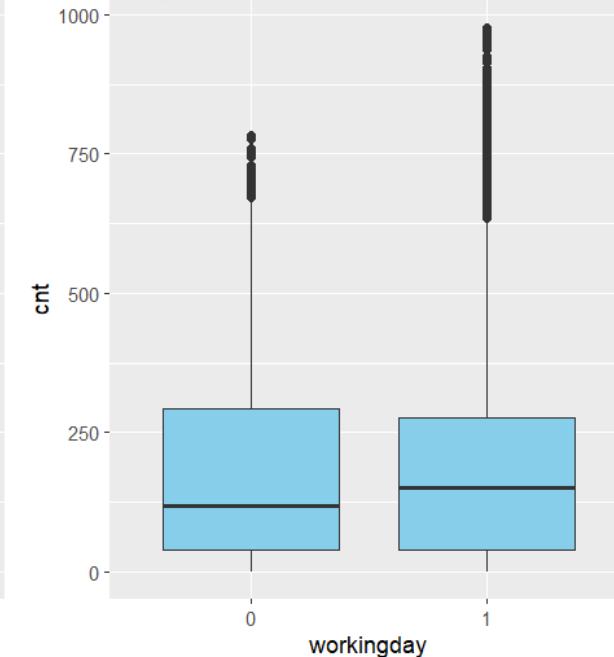
Box Plot for holiday



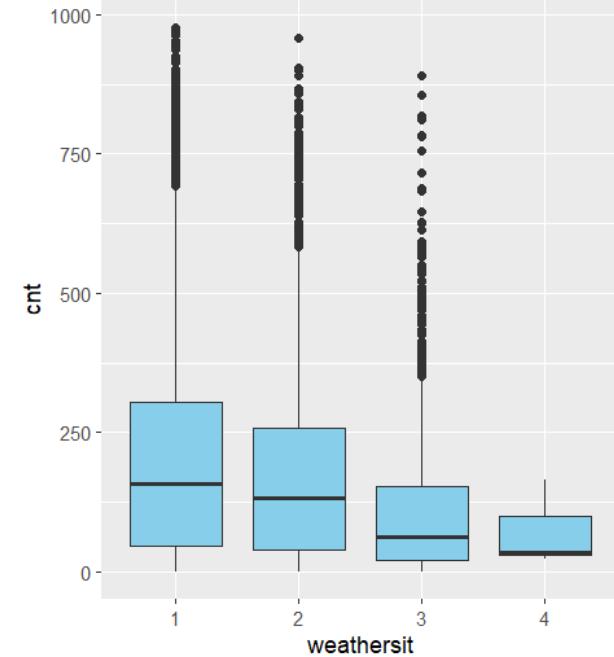
Box Plot for weekday



Box Plot for workingday

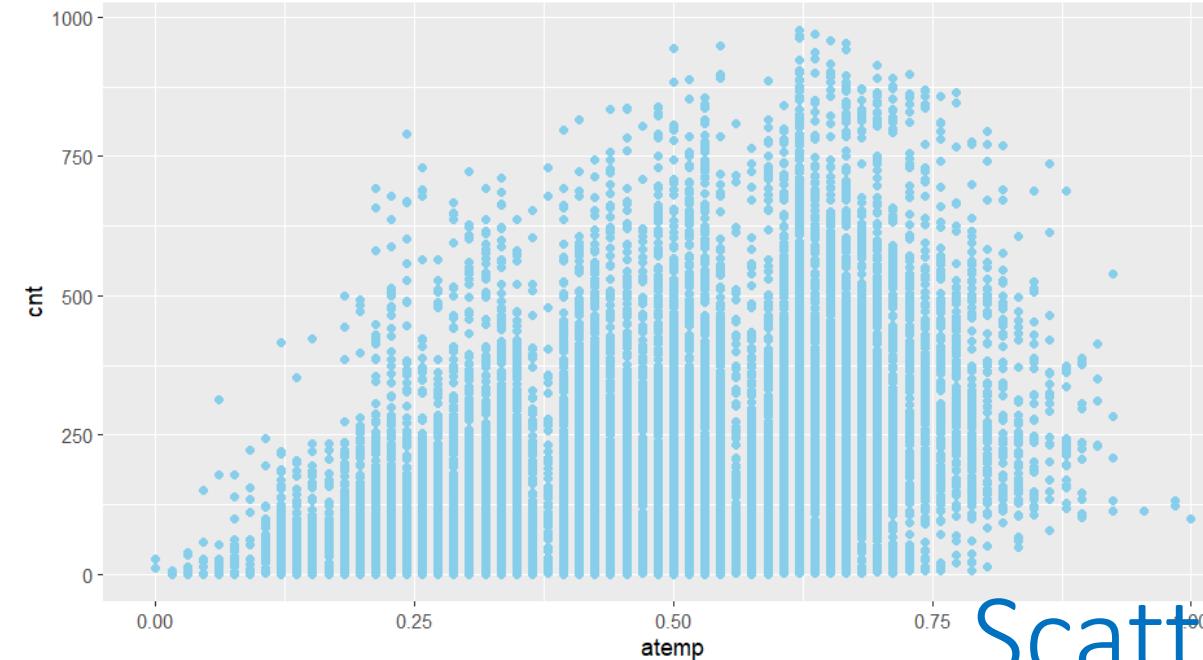


Box Plot for weathersit

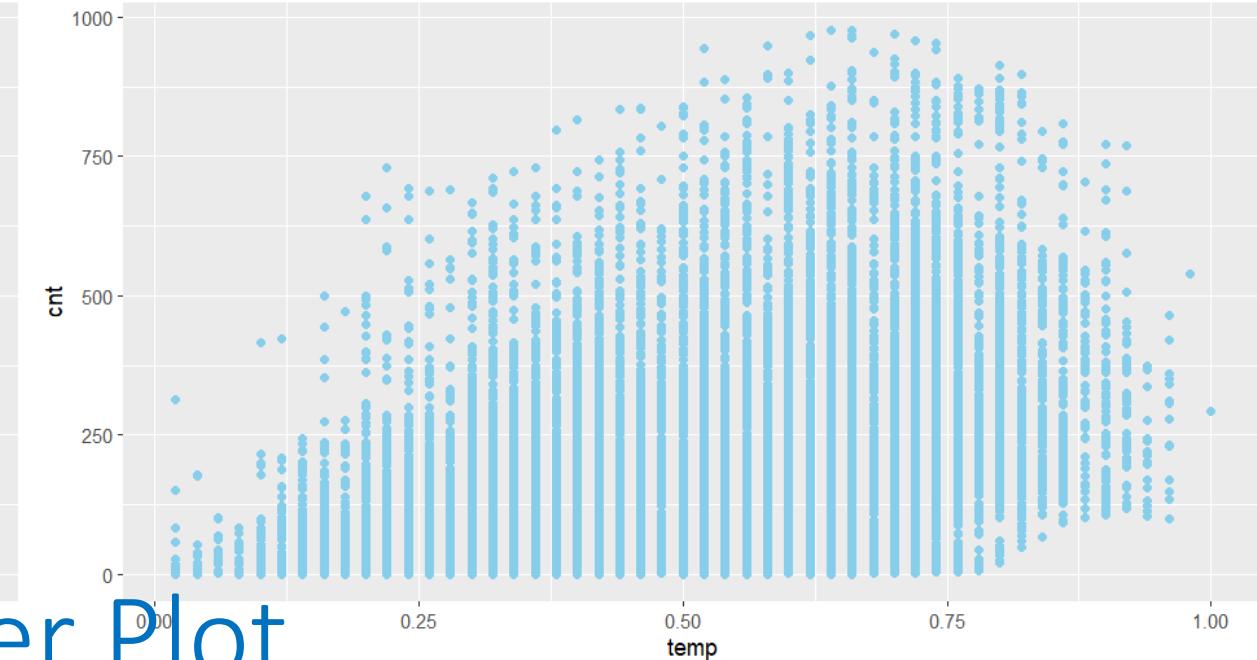


Box Plot

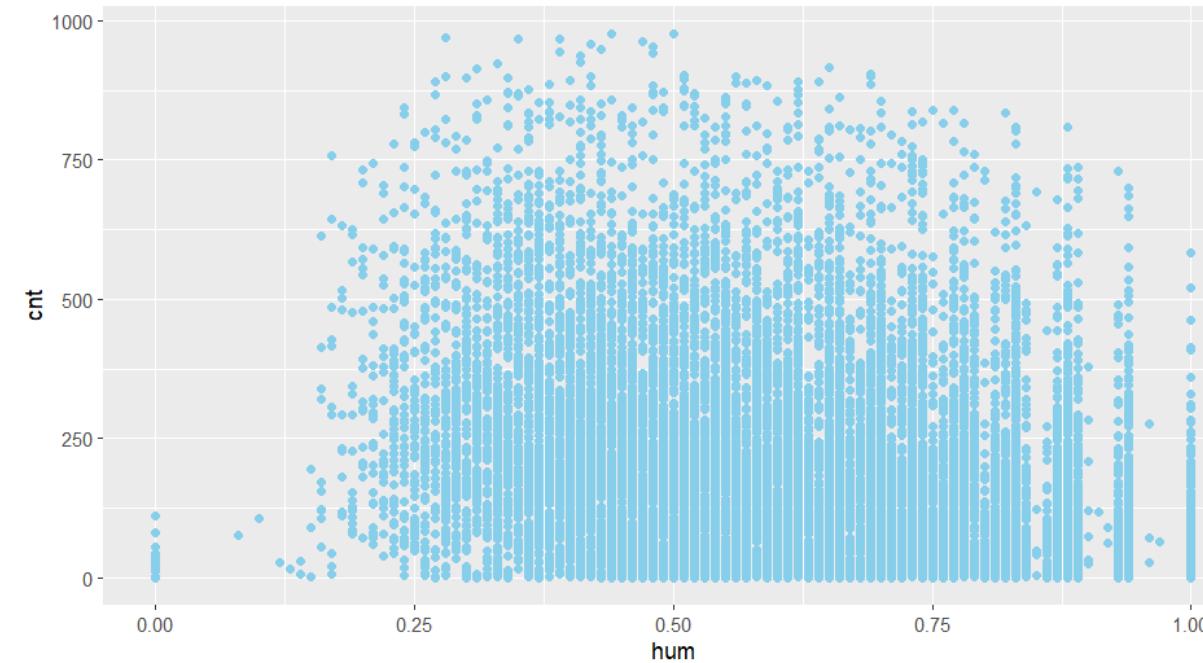
Scatter Plot for atemp vs cnt



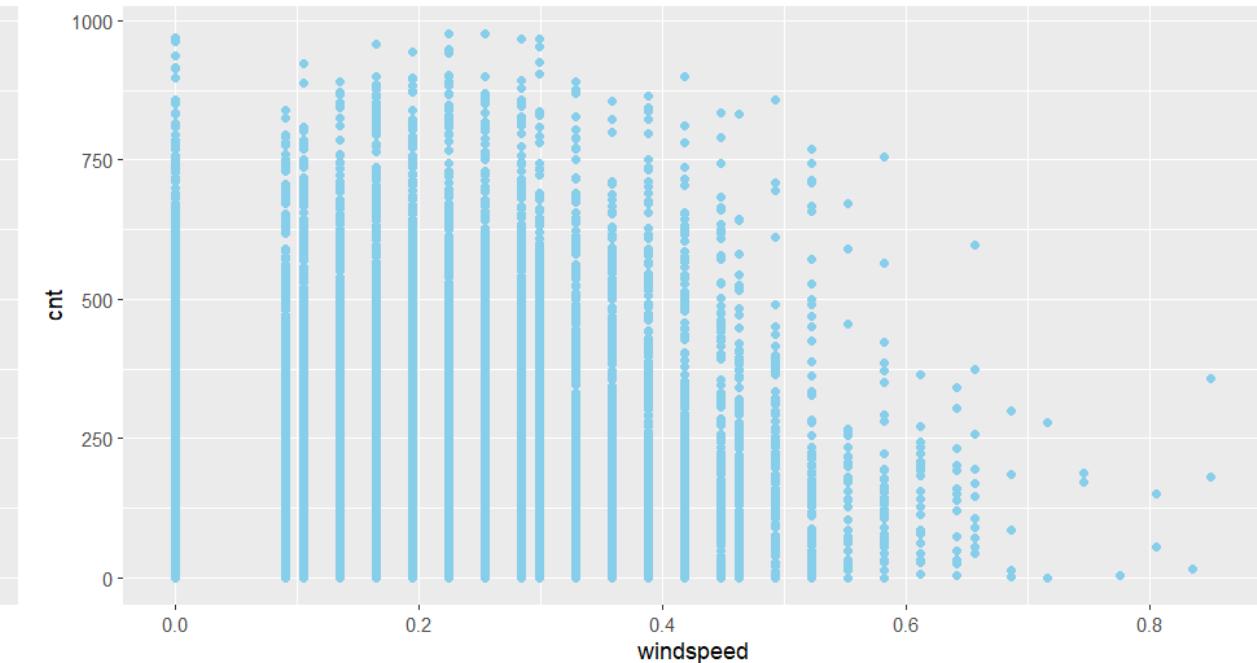
Scatter Plot for temp vs cnt



Scatter Plot for hum vs cnt

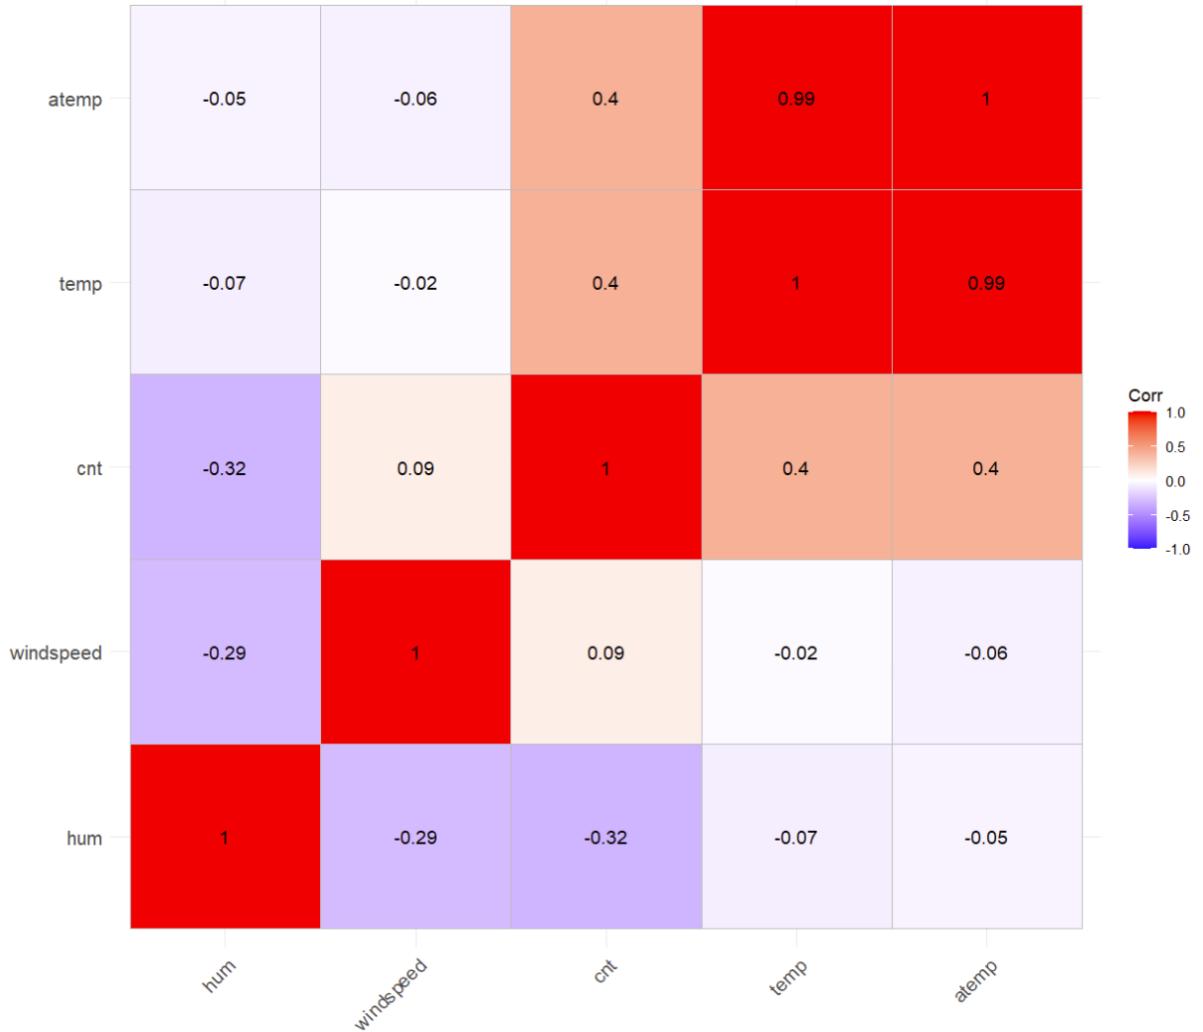


Scatter Plot for windspeed vs cnt

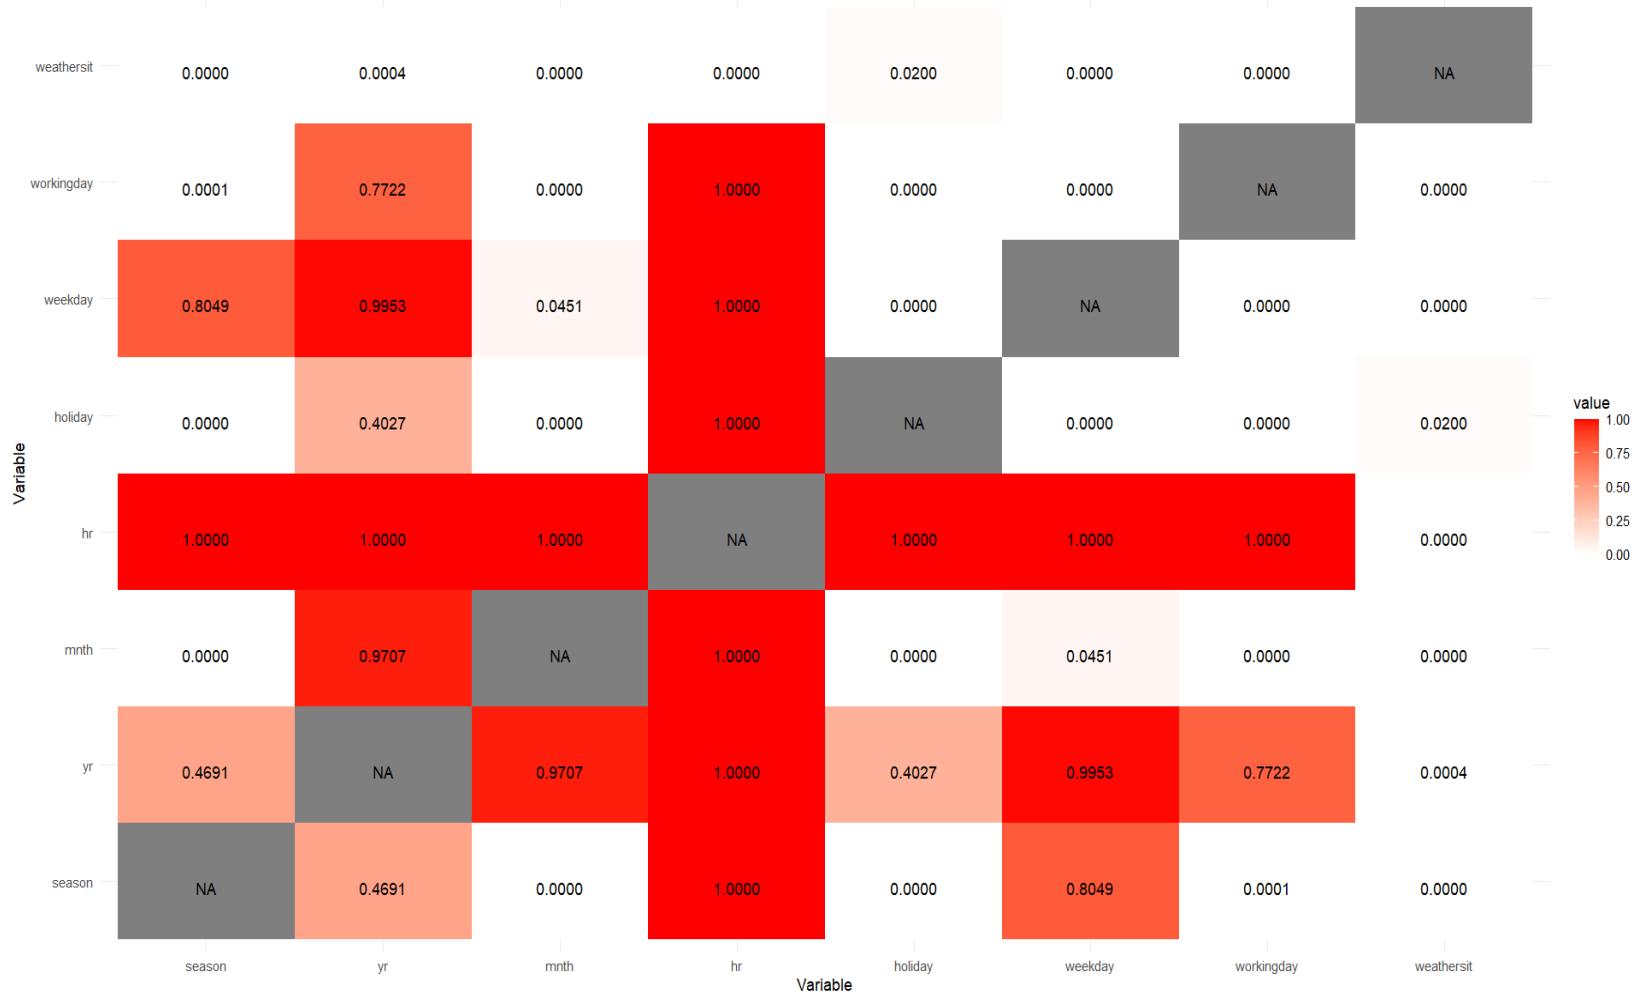


Scatter Plot

CORRELATION



CHI – SQUARE TEST



Transformation

```

> dim(dff)
[1] 17377   53
> colnames(dff)
[1] "season2"      "season3"      "season4"      "yr1"        "mnth2"       "mnth3"       "mnth4"       "mnth5"
[9] "mnth6"        "mnth7"        "mnth8"        "mnth9"       "mnth10"      "mnth11"      "mnth12"      "hr1"
[17] "hr2"          "hr3"          "hr4"          "hr5"         "hr6"         "hr7"         "hr8"         "hr9"
[25] "hr10"         "hr11"         "hr12"         "hr13"        "hr14"        "hr15"        "hr16"        "hr17"
[33] "hr18"         "hr19"         "hr20"         "hr21"        "hr22"        "hr23"        "holiday1"    "weekday1"
[41] "weekday2"     "weekday3"     "weekday4"     "weekday5"    "weekday6"    "weathersit2" "weathersit3" "weathersit4"
[49] "temp"          "atemp"        "hum"          "windspeed"   "cnt"         >

```

- Scaling using the Z normalization
 - ✓ Improved convergence for gradient Descent
 - ✓ Scale Invariance for some algorithmns
 - ✓ Interpretability and comparisons
 - ✓ Regularization Stability

- Identified and converted the categorical into factors and then converted into numerical using dummies.
- Scaled the numerical variable using the z Transformation.

weekday2	weekday3	weekday4	weekday5	weekday6	weathersit2	weathersit3	weathersit4	temp	atemp	hum	windspeed	cnt
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-1.3346564	-1.0932844	0.94738849	-1.55389398	16
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-1.4385237	-1.1817340	0.89555749	-1.55389398	40
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-1.4385237	-1.1817340	0.89555749	-1.55389398	32
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-1.3346564	-1.0932844	0.63640249	-1.55389398	13
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-1.3346564	-1.0932844	0.63640249	-1.55389398	1
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	1.6804765	-0.2981875	-0.01314009	-1.3346564	-1.2696017	0.63640249	-0.82152780	1
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-1.4385237	-1.1817340	0.89555749	-1.55389398	2
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-1.5423910	-1.2696017	1.20654350	-1.55389398	3
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-1.3346564	-1.0932844	0.63640249	-1.55389398	8
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-0.9191871	-0.7406499	0.68823349	-1.55389398	14
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-0.6075852	-0.4764649	0.68823349	0.51978124	36
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-0.7114525	-0.8290995	0.94738849	0.76417576	56
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	-0.2981875	-0.01314009	-0.3998506	-0.3001476	0.74006449	0.76417576	84
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	1.6804765	-0.2981875	-0.01314009	-0.1921160	-0.1238303	0.48090948	0.88596433	94
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	1.6804765	-0.2981875	-0.01314009	-0.1921160	-0.1238303	0.48090948	0.76417576	106
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	1.6804765	-0.2981875	-0.01314009	-0.2959833	-0.2116980	0.74006449	0.88596433	110
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	1.6804765	-0.2981875	-0.01314009	-0.3998506	-0.3001476	0.99921949	0.88596433	93
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	1.6804765	-0.2981875	-0.01314009	-0.2959833	-0.2116980	0.99921949	0.76417576	67
-0.405313	-0.4075237	-0.4071396	-0.408675	2.4331070	-0.5950351	3.3534014	-0.01314009	-0.3998506	-0.3001476	1.31020550	0.51978124	35

Feature Selection

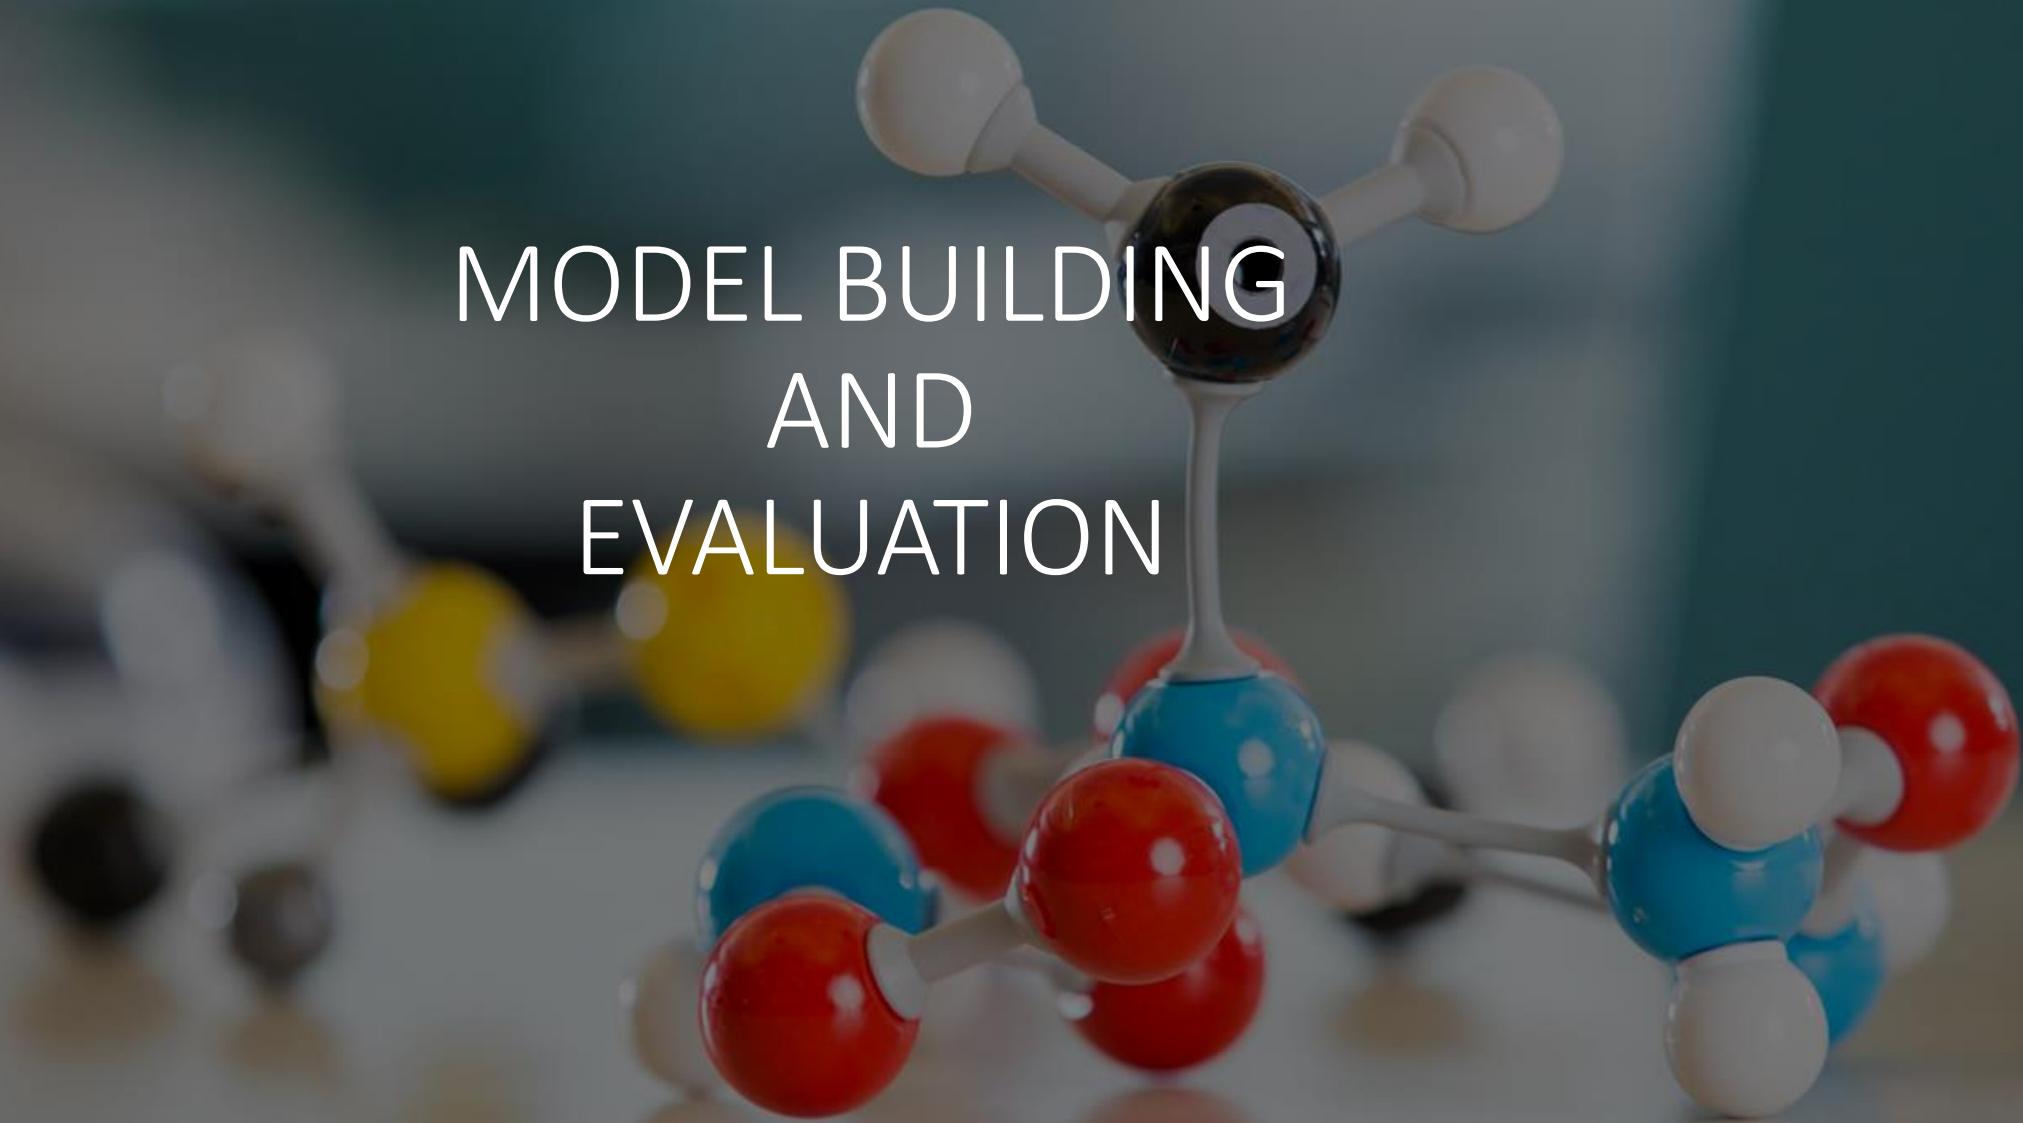
- “dteday” is a combination of “yr” and “mnth”.
- “instant” is index number
- “workingday” is associated with “holiday” and “weekday”
- “cnt” is also associated with “casual” and “registered”
- So, removed “dteday”, “instant”, “workingday”, “casual” and “registered.”

```
> colnames(data)
[1] "instant"    "dteday"     "season"      "yr"        "mnth"       "hr"        "holiday"    "weekday"   "registered"
[9] "workingday"  "weathersit"  "temp"       "atemp"     "hum"        "windspeed"  "casual"    "registered"
[17] "cnt"

> colnames(cleaned_data)
[1] "season"      "yr"         "mnth"       "hr"        "holiday"    "weekday"   "workingday" "weathersit"
[9] "temp"        "atemp"     "hum"        "windspeed"  "cnt"

> colnames(dff)
[1] "season2"    "season3"    "season4"    "yr1"       "mnth2"      "mnth3"      "mnth4"      "mnth5"
[9] "mnth6"       "mnth7"      "mnth8"      "mnth9"     "mnth10"     "mnth11"     "mnth12"     "hr1"
[17] "hr2"        "hr3"        "hr4"        "hr5"       "hr6"        "hr7"        "hr8"        "hr9"
[25] "hr10"       "hr11"       "hr12"       "hr13"     "hr14"       "hr15"       "hr16"       "hr17"
[33] "hr18"       "hr19"       "hr20"       "hr21"     "hr22"       "hr23"       "holiday1"   "weekday1"
[41] "weekday2"   "weekday3"   "weekday4"   "weekday5"  "weekday6"   "weathersit2" "weathersit3" "weathersit4"
[49] "temp"        "atemp"     "hum"        "windspeed" "cnt"
```

MODEL BUILDING AND EVALUATION



SIMPLE LINEAR REGRESSION



SIMPLE LINEAR REGRESSION

```
> lm.summary
```

Call:
lm(formula = .outcome ~ ., data = dat)

Residuals:

Min	1Q	Median	3Q	Max
-398.75	-60.55	-7.77	51.05	434.65

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	189.3061	0.8689	217.856	< 2e-16 ***
season2	16.4860	2.3838	6.916	4.86e-12 ***
season3	13.5262	2.8407	4.762	1.94e-06 ***
season4	29.8907	2.3664	12.632	< 2e-16 ***
yr1	42.2000	0.8798	47.967	< 2e-16 ***
mnth2	1.3089	1.1856	1.104	0.269622
mnth3	4.5463	1.3910	3.268	0.001084 **
mnth4	2.3462	2.0319	1.155	0.248238
mnth5	5.9408	2.2091	2.689	0.007169 **
mnth6	2.0166	2.2367	0.902	0.367290
mnth7	-2.5806	2.5506	-1.012	0.311672
mnth8	3.0686	2.4793	1.238	0.215848
mnth9	9.5621	2.1750	4.396	1.11e-05 ***
weathersit2	-4.7862	0.9498	-5.039	4.73e-07 ***
weathersit3	-18.6500	1.0008	-18.635	< 2e-16 ***
weathersit4	-0.8511	0.7794	-1.092	0.274861
temp	25.5675	6.1963	4.126	3.71e-05 ***
atemp	18.0229	5.7242	3.149	0.001644 **
hum	-15.2839	1.2053	-12.681	< 2e-16 ***
windspeed	-2.6369	0.9762	-2.701	0.006918 **

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 102.4 on 13849 degrees of freedom

Multiple R-squared: 0.6844, Adjusted R-squared: 0.6833

F-statistic: 577.6 on 52 and 13849 DF, p-value: < 2.2e-16

```
> lm.fit
```

Linear Regression

13902 samples

52 predictor

No pre-processing

Resampling: Cross-Validated (5 fold)

Summary of sample sizes: 11121, 11122, 11121, 11122, 11122

Resampling results:

RMSE	Rsquared	MAE
102.6258	0.6821557	75.82325

```
> cat("R-squared: ", r_squared_lm, "\n")
```

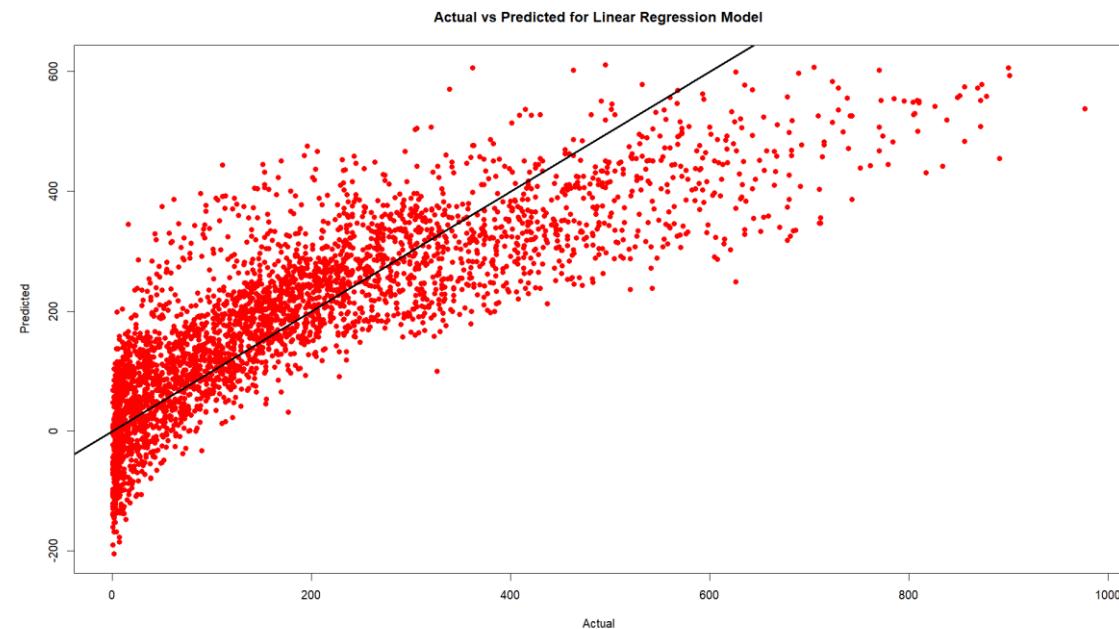
R-squared: 0.6929337

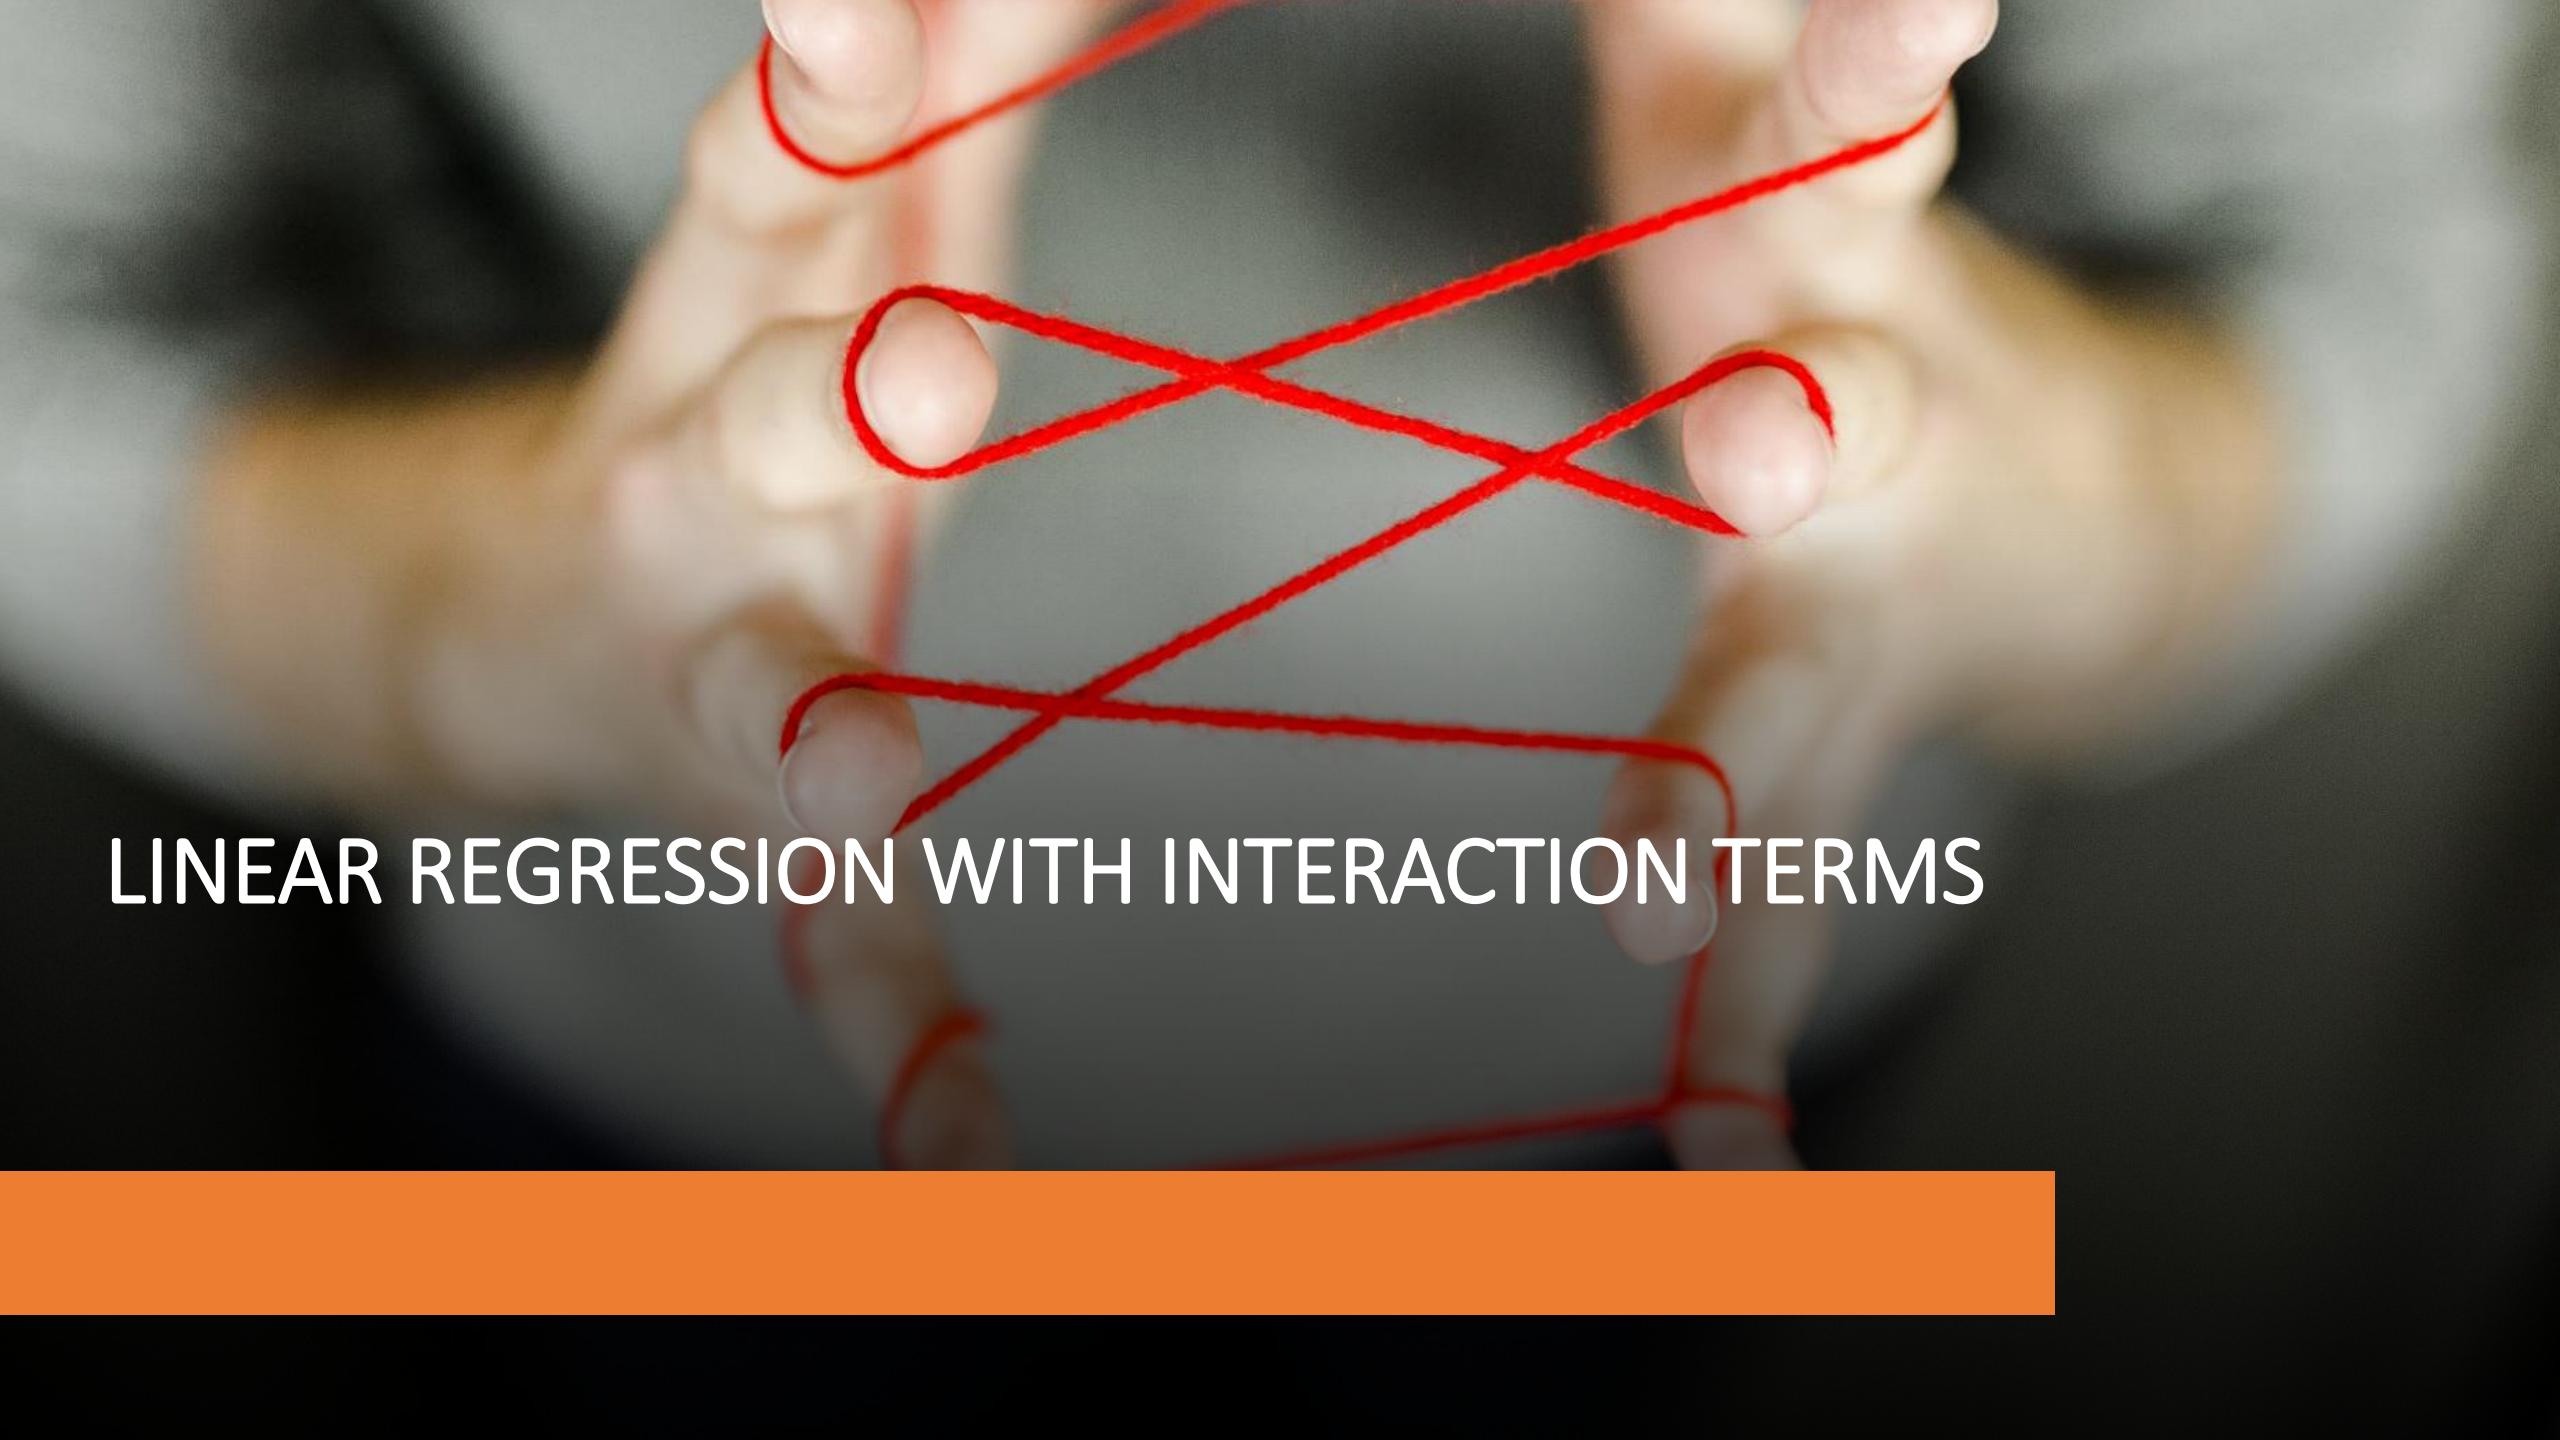
```
> cat("Mean Square Error (MSE): ", mse_lm, "\n")
```

Mean Square Error (MSE): 9841.131

```
> cat("Root Mean Square Error (RMSE): ", rmse_lm, "\n")
```

Root Mean Square Error (RMSE): 99.20247

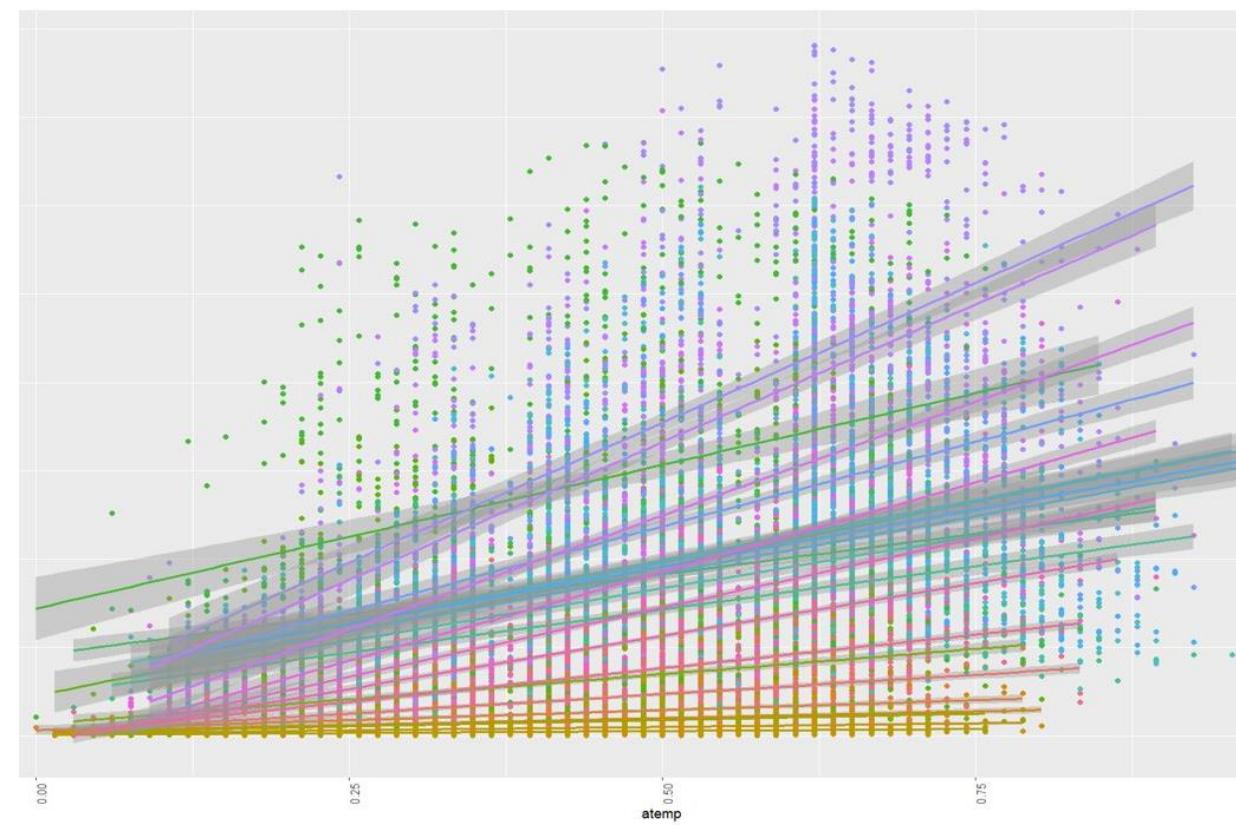
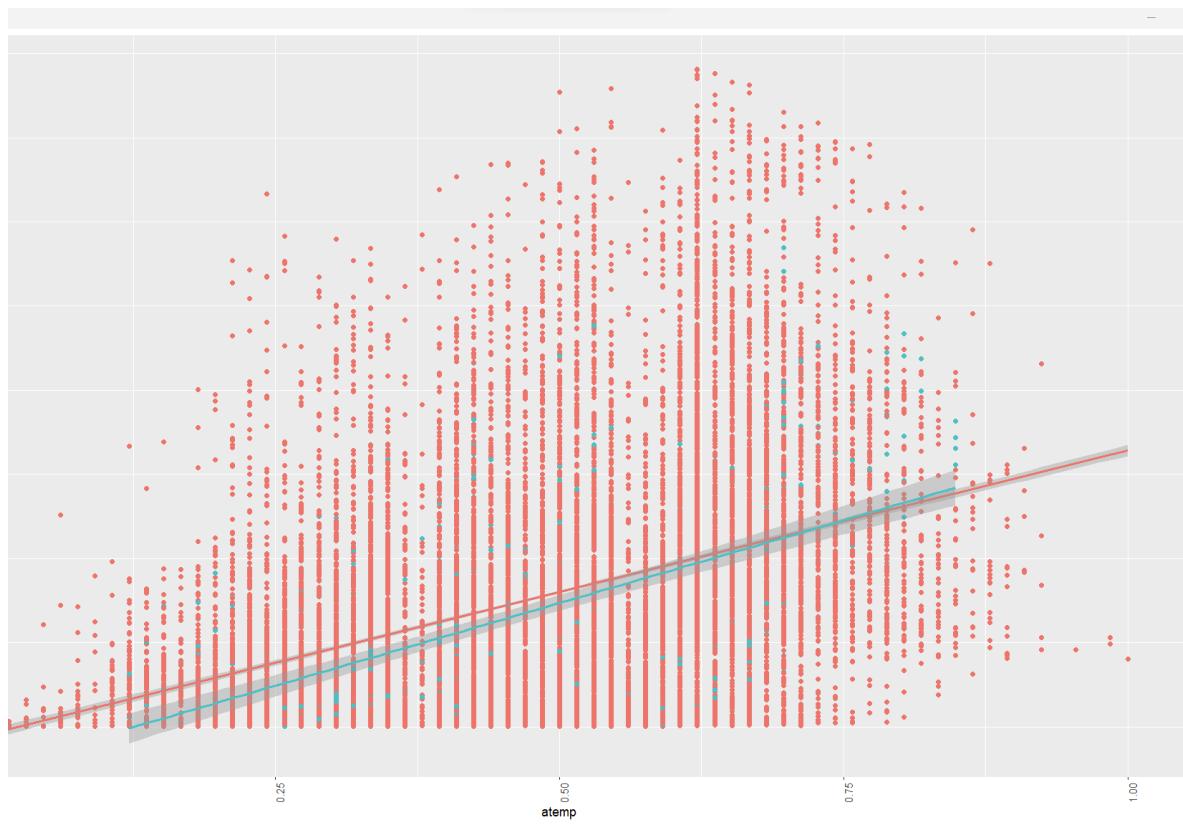


A close-up photograph of a person's hands holding a piece of red string. The string is knotted in a figure-eight pattern, with the loops extending upwards and downwards. The hands are positioned in the upper half of the frame, with fingers gripping the ends of the string.

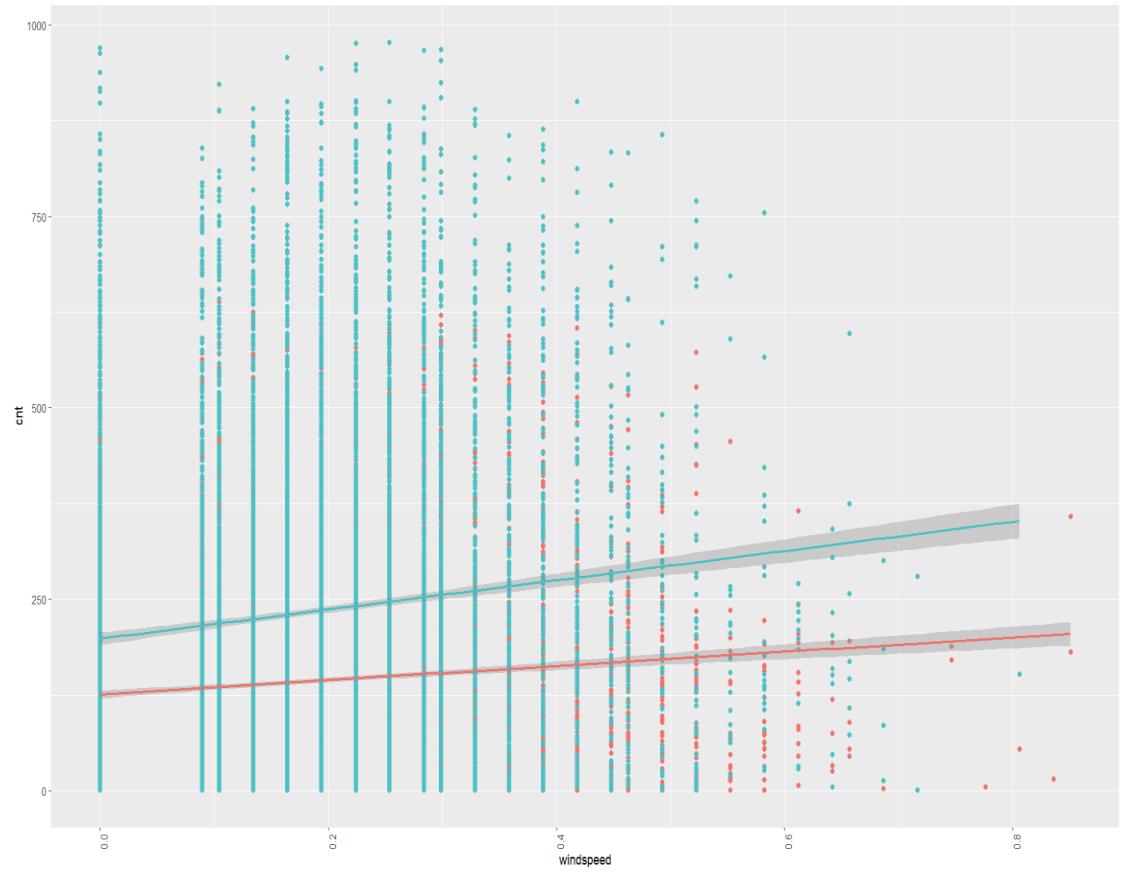
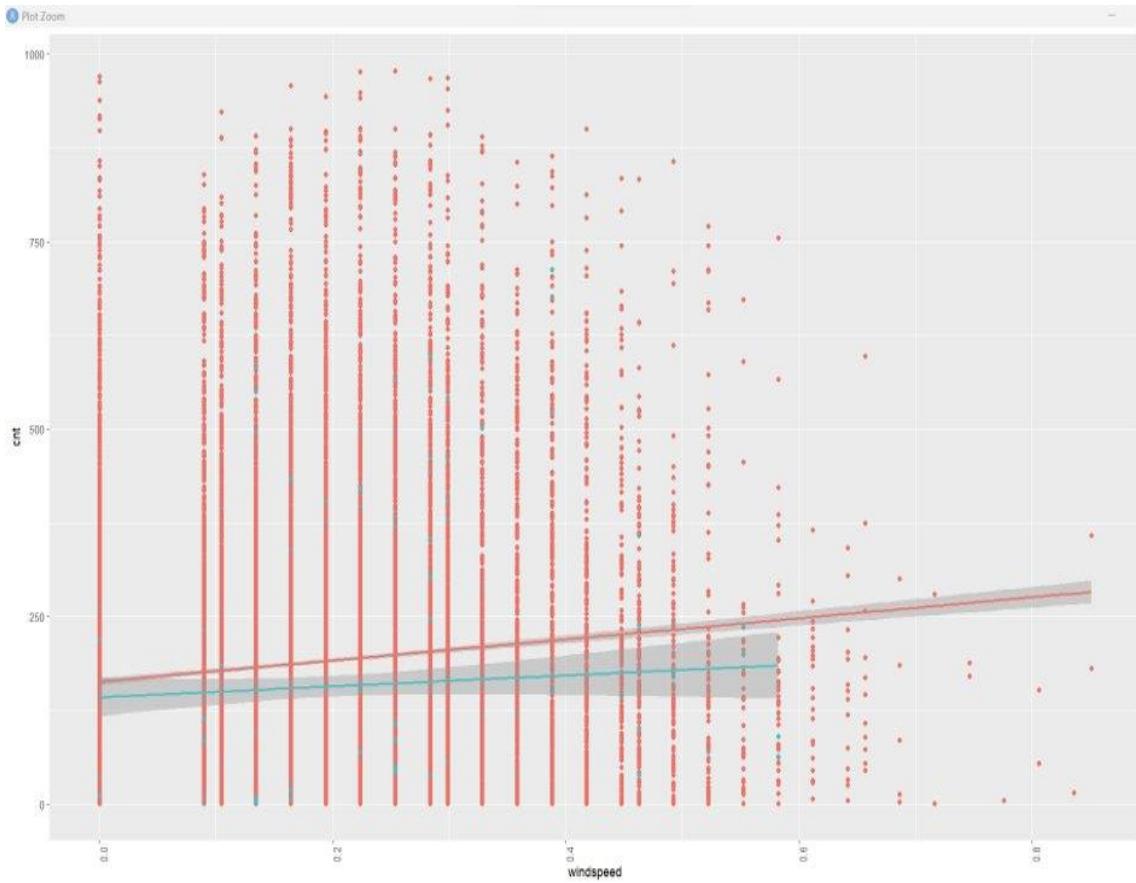
LINEAR REGRESSION WITH INTERACTION TERMS

LINEAR REGRESSION WITH INTERACTION TERMS

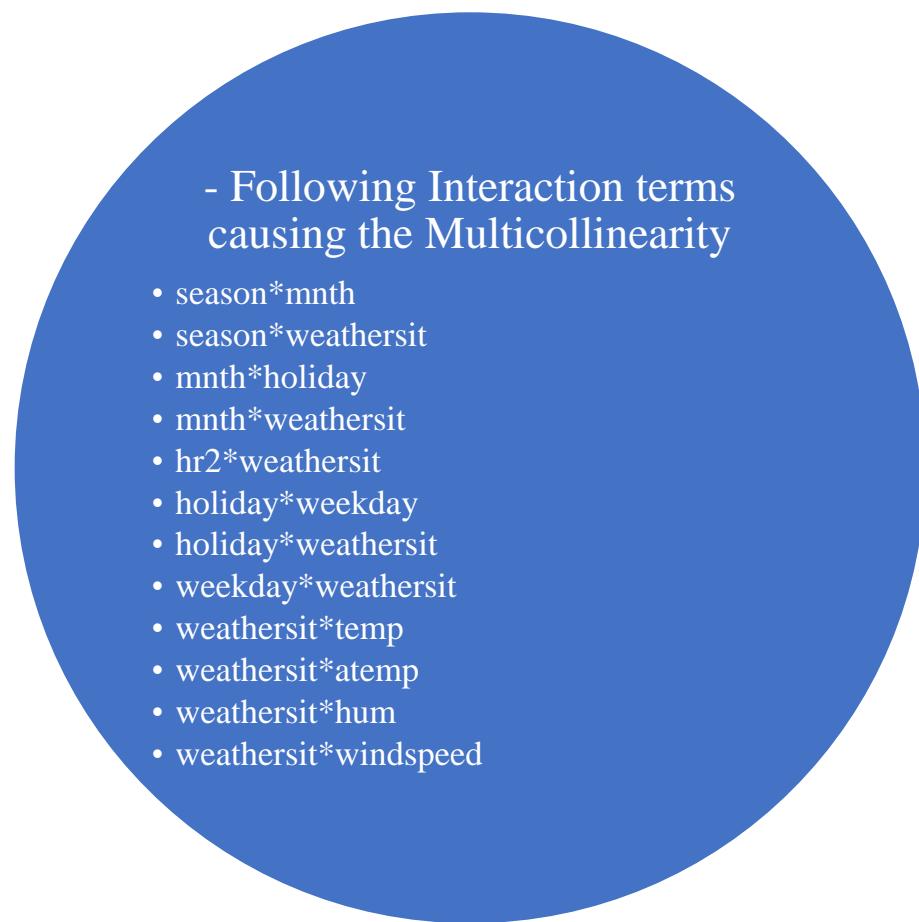
INTERACTION TERM



NON - INTERACTION TERM



LINEAR REGRESSION WITH INTERACTION TERMS



- Finally considered the Interaction terms

- ✓ holiday*atemp
- ✓ temp*atemp
- ✓ season*yr
- ✓ yr*atemp
- ✓ hr*holiday
- ✓ season*hr
- ✓ yr*hr
- ✓ mnth*hr
- ✓ yr*hr
- ✓ mnth*hr
- ✓ yr*hr
- ✓ mnth*hum
- ✓ yr*mnth
- ✓ hr*weekday

LINEAR REGRESSION WITH INTERACTION TERMS

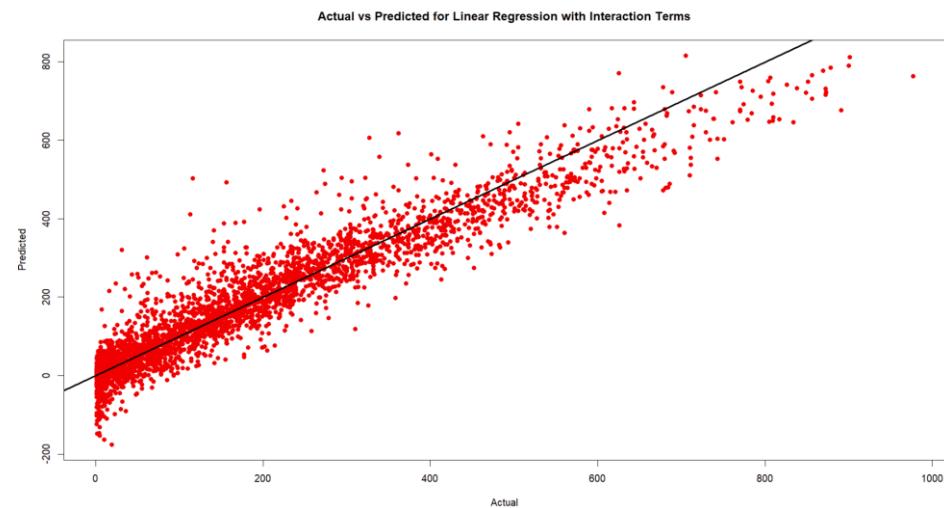
```
> lm.fit_it <- train(cnt ~.+(holiday*atemp)+(temp*atemp)+(season*yr)+(yr*atemp)+(hr*holiday)+(season*hr)+  
+                               (yr*hr)+(mnth*hr)+(yr*hr)+(mnth*hr)+(yr*hr)+(mnth*hum)+(yr*mnth)+(hr*weekday),  
+                               data = cleaned_data[trainIndex, ],method="lm",trControl = ctrl)  
> lm.summary_it  
> lm.summary_it  
  
Call:  
lm(formula = .outcome ~ ., data = dat)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-471.46 -24.72 -0.66  27.37 458.02  
  
Coefficients:  
            Estimate Std. Error t value Pr(>|t|)  
(Intercept) 25.4113   12.4849   2.035 0.041835 *  
season2     -6.8703   14.7611  -0.465 0.641627  
season3     -5.4973   18.4210  -0.298 0.765385  
season4     -1.6012   16.1243  -0.099 0.920897  
yr1        -47.1773   6.6346  -7.111 1.21e-12 ***  
mnth2     -25.7637   14.8869  -1.731 0.083542 .  
mnth3     -43.4961   15.4770  -2.810 0.004956 **  
mnth4      12.1285   21.4464   0.566 0.571724  
mnth5      66.0083   22.5466   2.928 0.003421 **  
mnth6      21.5619   21.8351   0.987 0.323421  
mnth7     -28.5266   25.2783  -1.129 0.259129  
mnth8      47.7246   25.3290   1.884 0.059562 .  
mnth9     141.9655   24.9531   5.689 1.30e-08 ***  
mnth10     96.3250   23.4073   4.115 3.89e-05 ***  
  
.`mnth11:hr2` -11.5817   27.6573  -0.419 0.675400  
.`mnth12:hr2` -13.1250   22.3477  -0.587 0.557004  
.`mnth2:hr3`  -4.2622   17.7005  -0.241 0.809718  
.`mnth3:hr3`  -7.6630   18.4342  -0.416 0.677642  
.`mnth4:hr3`  -0.5897   26.7212  -0.022 0.982392  
.`mnth5:hr3` -25.5080   26.5778  -0.960 0.337200  
[ reached getOption("max.print") -- omitted 387 rows ]  
---  
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
Residual standard error: 57.79 on 13315 degrees of freedom  
Multiple R-squared: 0.9034, Adjusted R-squared: 0.8991  
F-statistic: 212.5 on 586 and 13315 DF, p-value: < 2.2e-16
```

MODEL

```
> lm.fit_it  
Linear Regression  
13902 samples  
11 predictor  
No pre-processing  
Resampling: Cross-validated (5 fold)  
Summary of sample sizes: 11120, 11122, 11122, 11122, 11122  
Resampling results:  
  RMSE     Rsquared     MAE  
  59.62651  0.8926881  40.74902
```

METRICS WITH TEST DATA

```
> cat("R-squared: ", r_squared_it, "\n")  
R-squared: 0.8994243  
> cat("Mean Square Error (MSE): ", mse_it, "\n")  
Mean Square Error (MSE): 3223.339  
> cat("Root Mean Square Error (RMSE): ", rmse_it, "\n")  
Root Mean Square Error (RMSE): 56.77445  
> |
```



LINEAR REGRESSION WITH POLYNOMIAL DEGREE



LINEAR REGRESSION WITH POLYNOMIAL DEGREE

Results for degree 2

```
> lm.fit_2
Linear Regression

13902 samples
 56 predictor

No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 11121, 11122, 11121, 11122, 11122
Resampling results:

RMSE      Rsquared     MAE
102.4031  0.6834022  75.65933
```

Results for degree 4

```
> lm.fit_4
Linear Regression

13902 samples
 64 predictor

No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 11122, 11121, 11122, 11122, 11121
Resampling results:

RMSE      Rsquared     MAE
101.0793  0.6915714  75.20426
```

Results for degree 3

```
> lm.fit_3
Linear Regression

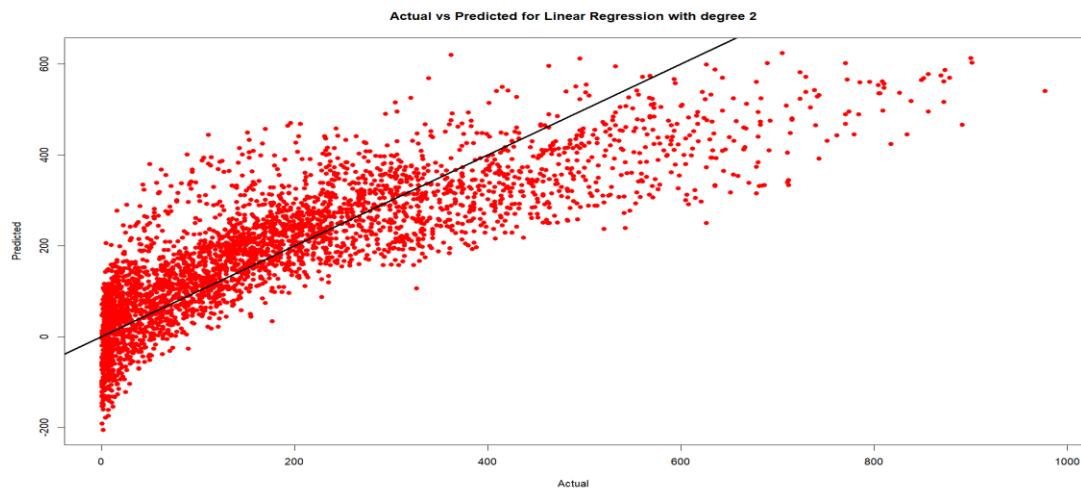
13902 samples
 60 predictor

No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 11122, 11121, 11121, 11122, 11122
Resampling results:

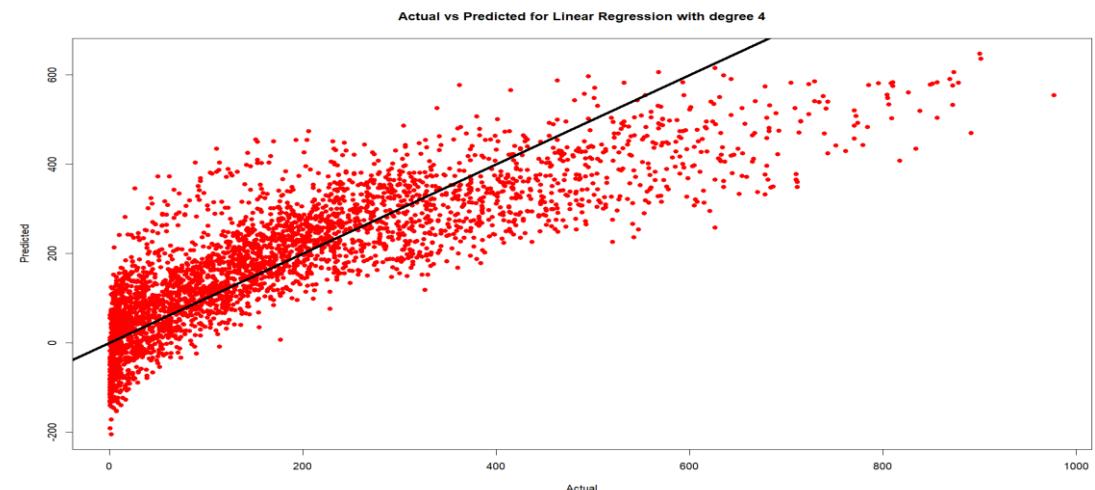
RMSE      Rsquared     MAE
101.2993  0.6901062  75.29893
```

LINEAR REGRESSION WITH POLYNOMIAL DEGREE

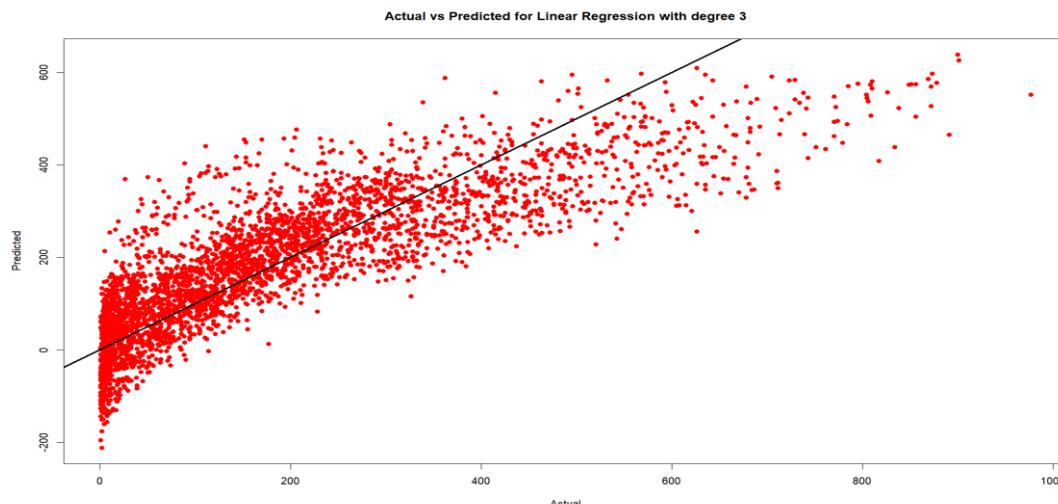
Results for degree 2



Results for degree 4



Results for degree 3



LINEAR REGRESSION WITH POLYNOMIAL DEGREE

Results for degree 2

```
> cat("R-squared: ", r_squared_lm2, "\n")
R-squared: 0.6954175
> cat("Mean Square Error (MSE): ", mse_lm2, "\n")
Mean Square Error (MSE): 9761.529
> cat("Root Mean Square Error (RMSE): ", rmse_lm2, "\n")
Root Mean Square Error (RMSE): 98.80045
```

Results for degree 4

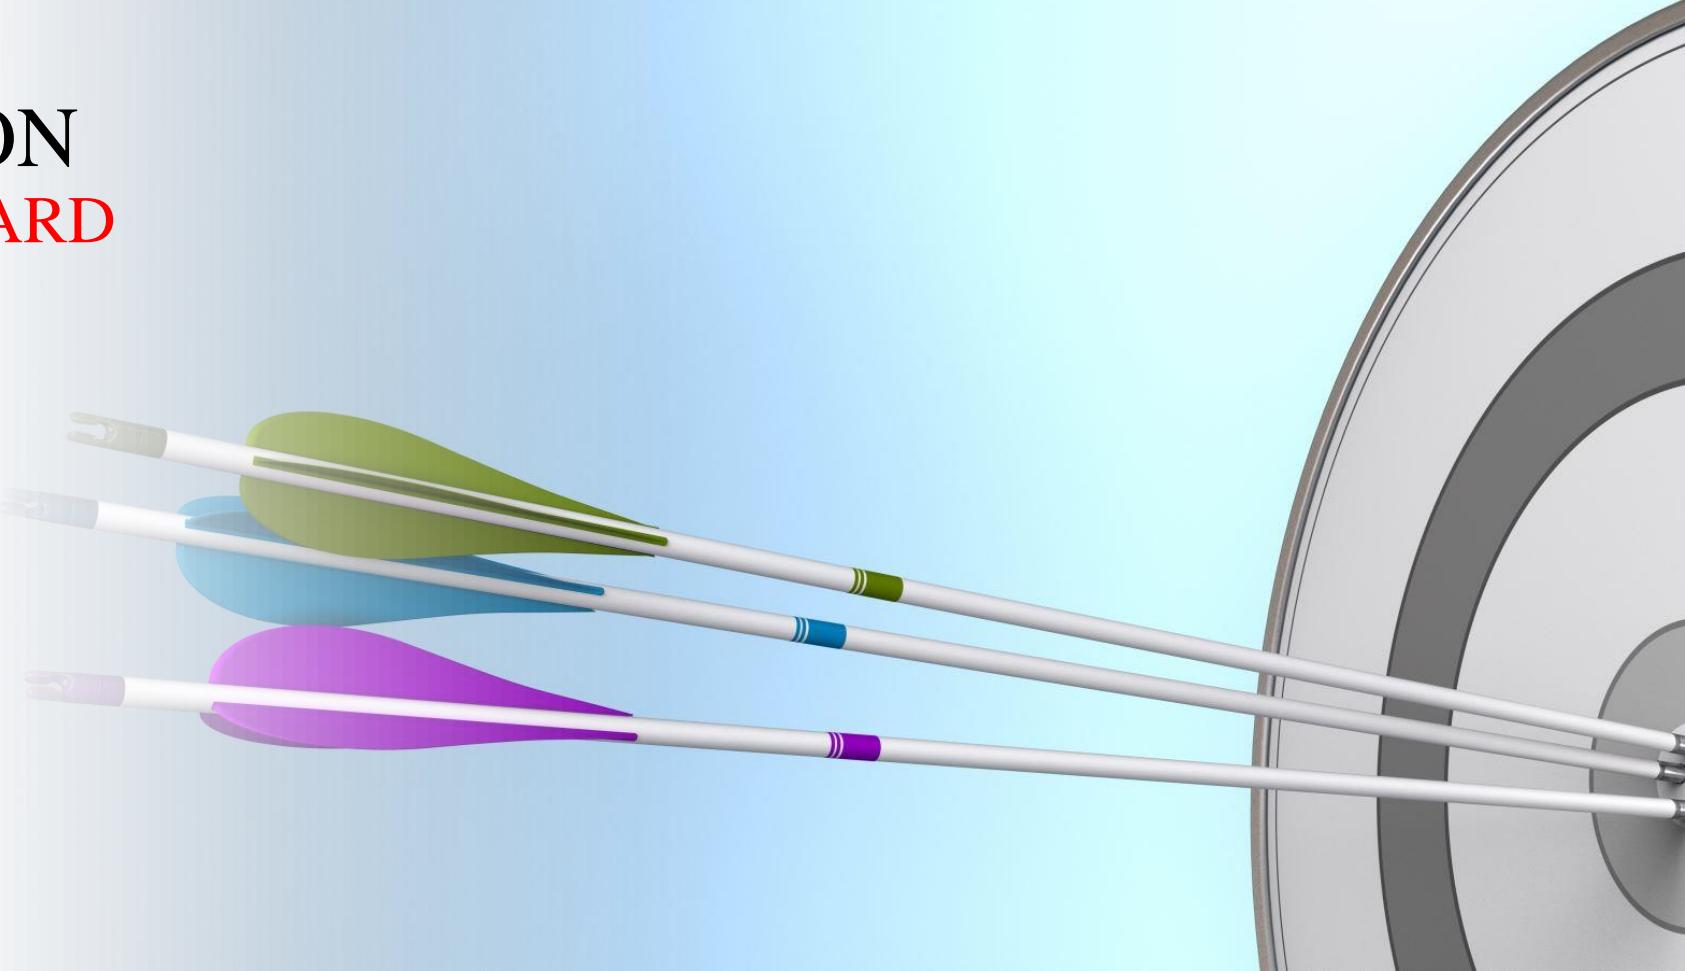
```
> cat("R-squared: ", r_squared_lm4, "\n")
R-squared: 0.7031851
> cat("Mean Square Error (MsE): ", mse_lm4, "\n")
Mean Square Error (MsE): 9512.586
> cat("Root Mean Square Error (RMSE): ", rmse_lm4, "\n")
Root Mean Square Error (RMSE): 97.53249
```

Results for degree 3

```
> cat("R-squared: ", r_squared_lm3, "\n")
R-squared: 0.7011467
> cat("Mean Square Error (MsE): ", mse_lm3, "\n")
Mean Square Error (MsE): 9577.911
> cat("Root Mean Square Error (RMSE): ", rmse_lm3, "\n")
Root Mean Square Error (RMSE): 97.8668
```

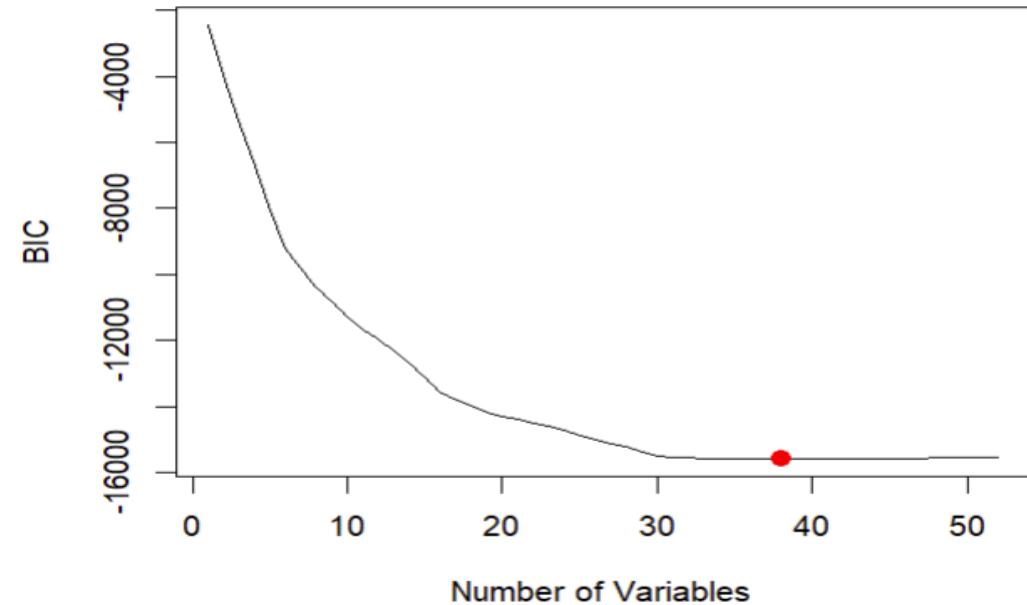
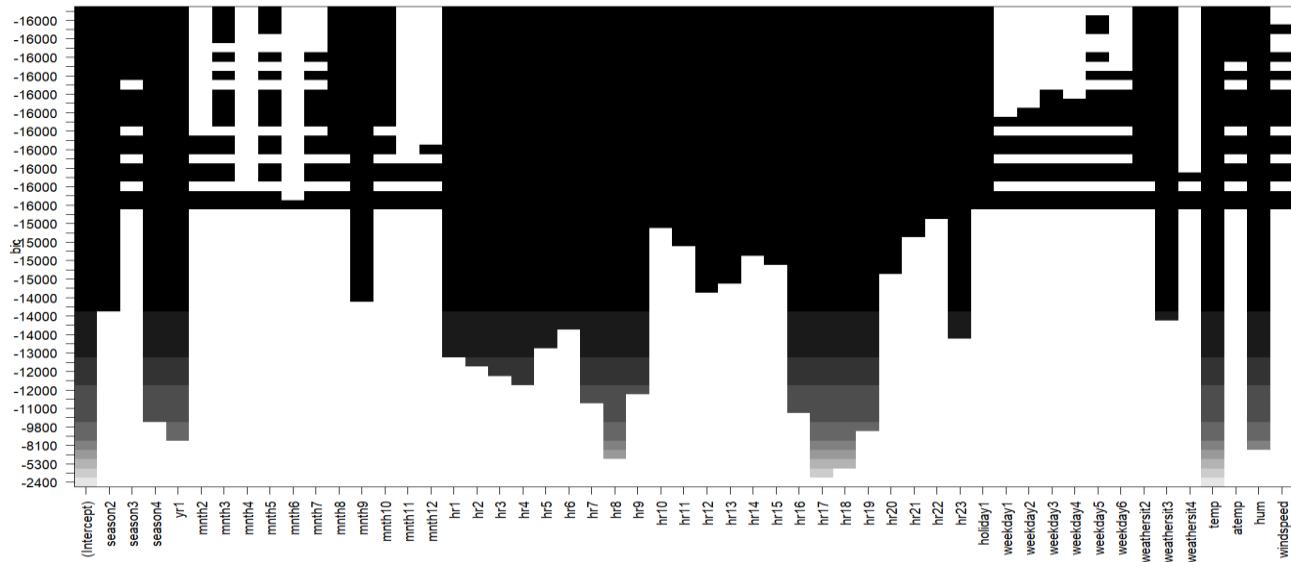
SUBSET SELECTION

FORWARD AND BACKWARD SELECTION



SUBSET SELECTION

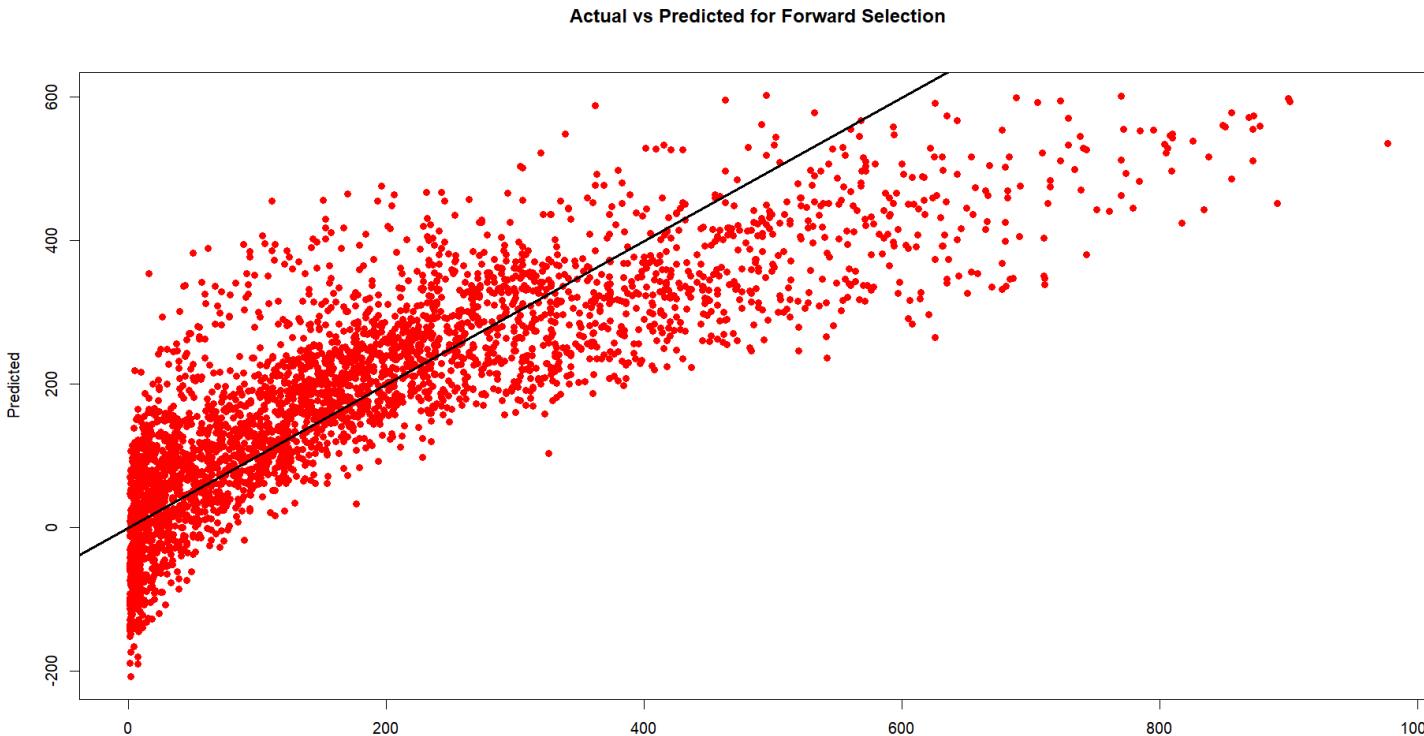
FORWARD SELECTION



```
> coef(regfit.fwd, 38)
```

(Intercept)	season2	season3	season4	yr1	mnth3	mnth5	mnth8	mnth9
189.288952	18.747791	10.223475	25.909799	42.103023	3.584787	3.924685	4.702713	11.367324
mnth10	hr1	hr2	hr3	hr4	hr5	hr6	hr7	hr8
5.544430	-3.209394	-5.243933	-6.922451	-8.122287	-4.202379	7.272126	34.115324	62.272467
hr9	hr10	hr11	hr12	hr13	hr14	hr15	hr16	hr17
32.447657	21.261054	26.891246	34.650584	33.425537	30.280654	32.611738	44.998609	76.803306
hr18	hr19	hr20	hr21	hr22	hr23	holiday1	weathersit2	weathersit3
68.946476	46.829041	31.178625	21.679811	14.347000	6.588060	-4.361285	-4.638782	-18.646796
temp	atemp	hum						
24.600633	19.841455	-15.254342						

FORWARD SELECTION

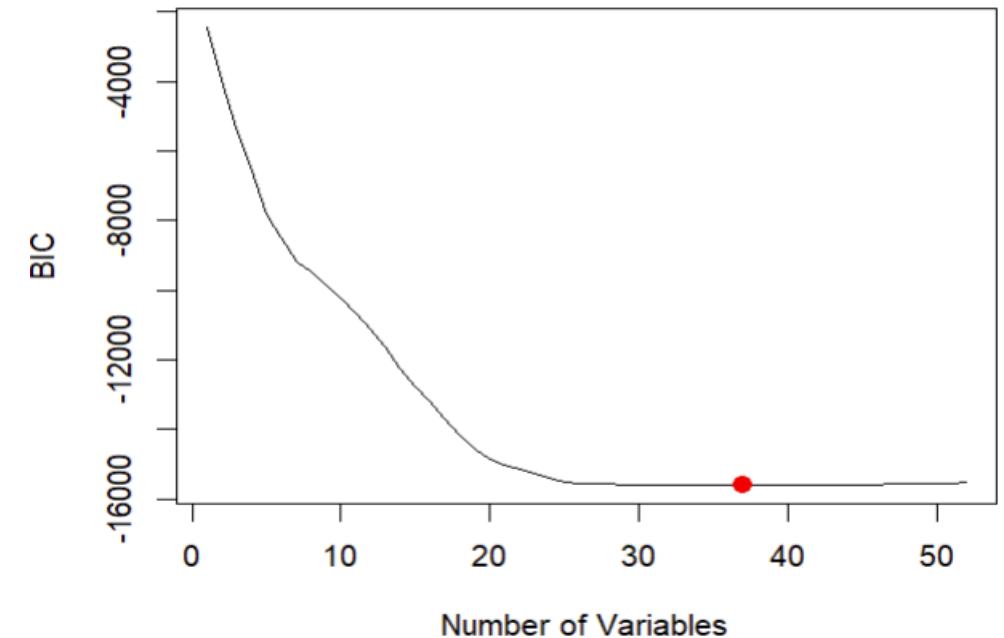
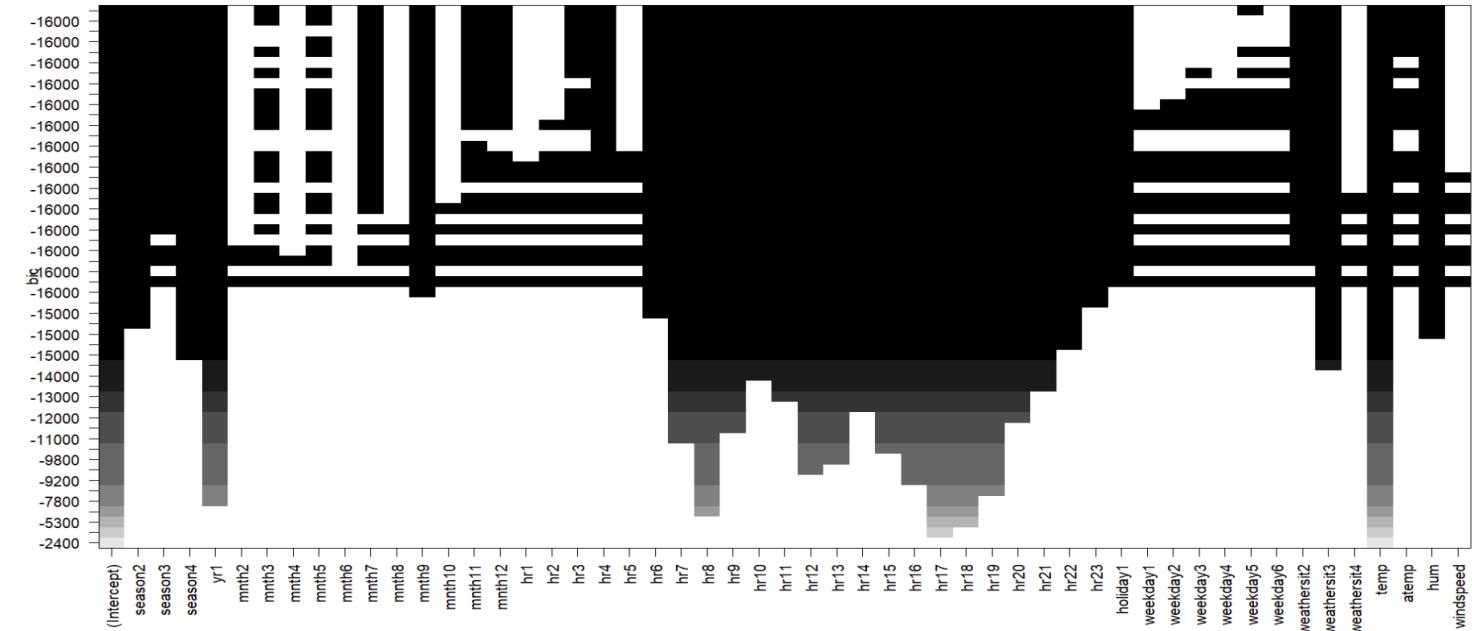


```
> cat("R-squared: ", r_squared_fw, "\n")
R-squared: 0.6909245
> cat("Mean Square Error (MsE): ", mse_fw, "\n")
Mean Square Error (MsE): 9905.524
> cat("Root Mean Square Error (RMSE): ", rmse_fw, "\n")
Root Mean Square Error (RMSE): 99.5265
```

```
> cv.errors.fwd
[1] 26444.648 24318.891 22022.468 20428.590 18292.142 16597.379 15734.361 15044.118 14638.365 14045.652 13643.904 13355.286 12989.082 12584.272
[15] 12214.863 11711.843 11519.840 11342.025 11218.695 11078.790 11015.072 10923.049 10823.611 10693.821 10637.727 10504.583 10406.582 10334.992
[29] 10174.868 10083.815 10045.140 10028.980 10016.478 9966.989 9956.881 9923.796 9925.677 9905.524 9901.974 9875.581 9858.089 9846.435
[43] 9848.432 9837.912 9837.510 9834.092 9837.029 9842.000 9839.138 9839.497 9840.846
> rmse_fwd
[1] 162.61811 155.94515 148.39969 142.92862 135.24845 128.83082 125.43668 122.65447 120.98911 118.51435 116.80712 115.56507 113.96965 112.17964
[15] 110.52087 108.22127 107.33052 106.49894 105.91834 105.25583 104.95271 104.51339 104.03658 103.41093 103.13936 102.49187 102.01265 101.66116
[29] 100.87055 100.41820 100.22545 100.14479 100.08235 99.83481 99.78417 99.61825 99.62769 99.52650 99.50866 99.37596 99.28791 99.22920
[43] 99.23927 99.18625 99.18422 99.16699 99.18180 99.20685 99.19243 99.19424 99.20104
> cv.r2.fwd
[1] 0.1748651 0.2411937 0.3128474 0.3625802 0.4292423 0.4821229 0.5090511 0.5305883 0.5432487 0.5617428 0.5742782 0.5832838 0.5947102 0.6073413
[15] 0.6188677 0.6345631 0.6405540 0.6461023 0.6499504 0.6543158 0.6563040 0.6591753 0.6622780 0.6663277 0.6680780 0.6722324 0.6752903 0.6775241
[29] 0.6825203 0.6853614 0.6865681 0.6870724 0.6874624 0.6890066 0.6893220 0.6903543 0.6902956 0.6909245 0.6910352 0.6918587 0.6924045 0.6927682
[43] 0.6927059 0.6930341 0.6930467 0.6931533 0.6930617 0.6929066 0.6929959 0.6929847 0.6929426
```

SUBSET SELECTION

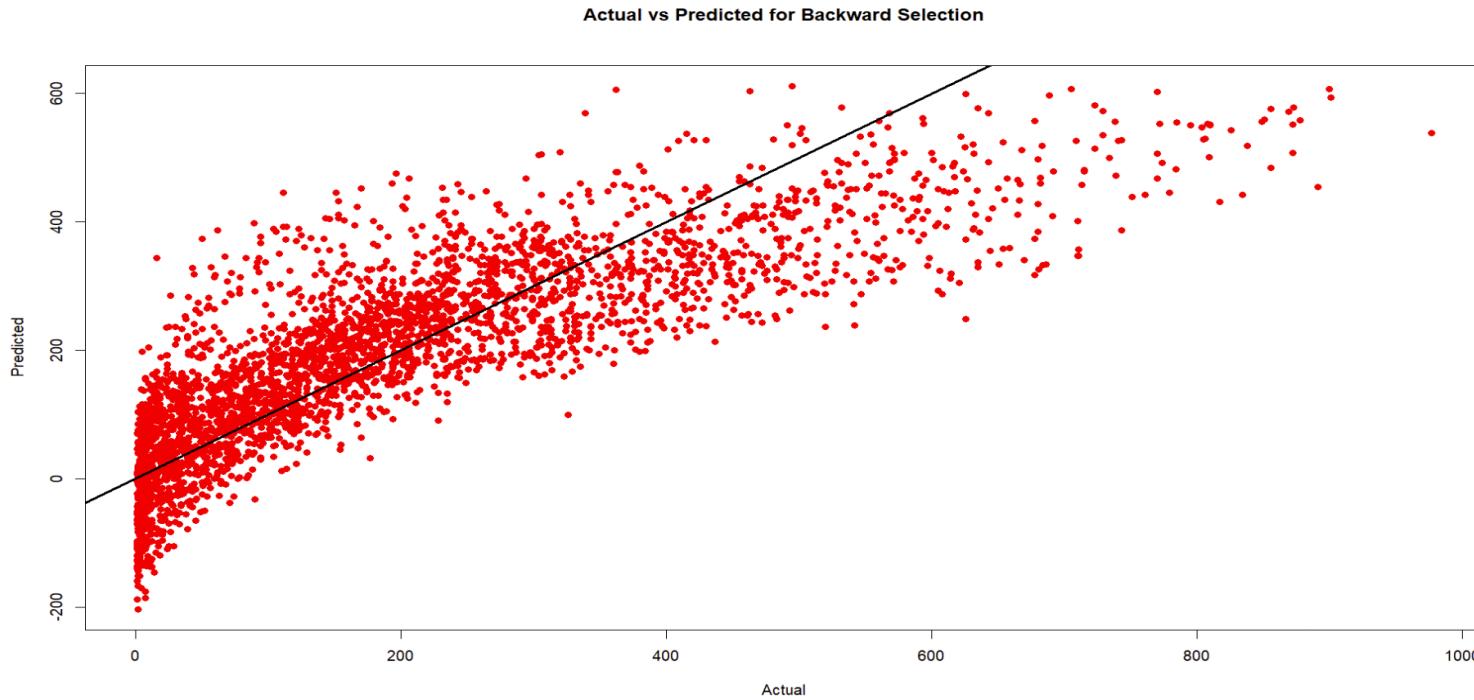
BACKWARD SELECTION



```
> coef(regfit.fwd, 38)
```

(Intercept)	season2	season3	season4	yr1	mnth3	mnth5	mnth8	mnth9
189.288952	18.747791	10.223475	25.909799	42.103023	3.584787	3.924685	4.702713	11.367324
mnth10	hr1	hr2	hr3	hr4	hr5	hr6	hr7	hr8
5.544430	-3.209394	-5.243933	-6.922451	-8.122287	-4.202379	7.272126	34.115324	62.272467
hr9	hr10	hr11	hr12	hr13	hr14	hr15	hr16	hr17
32.447657	21.261054	26.891246	34.650584	33.425537	30.280654	32.611738	44.998609	76.803306
hr18	hr19	hr20	hr21	hr22	hr23	holiday1	weathersit2	weathersit3
68.946476	46.829041	31.178625	21.679811	14.347000	6.588060	-4.361285	-4.638782	-18.646796
temp	atemp	hum						
24.600633	19.841455	-15.254342						

BACKWARD SELECTION



```
> cat("R-squared: ", r_squared_bw, "\n")
R-squared: 0.6906202
> cat("Mean Square Error (MsE): ", mse_bw, "\n")
Mean Square Error (MsE): 9915.277
> cat("Root Mean Square Error (RMSE): ", rmse_bw, "\n")
Root Mean Square Error (RMSE): 99.57548
~ |
```

```
> cv.errors.bwd
[1] 26444.648 24318.891 22022.468 20428.590 18488.851 17397.522 16699.877 16275.440 15902.588 15602.269 15092.865 14563.414 13947.860 13212.865
[15] 12677.919 12256.154 11730.837 11316.816 10918.783 10667.190 10517.273 10385.000 10290.135 10184.860 10119.078 10081.399 10064.493 10056.063
[29] 10018.403 10018.208 9991.928 9986.486 9970.173 9934.566 9918.827 9918.969 9915.277 9903.889 9905.291 9893.876 9892.542 9890.012
[43] 9888.841 9878.748 9869.655 9844.325 9844.702 9836.636 9838.494 9839.497 9840.846
> rmse_bwd
[1] 162.61811 155.94515 148.39969 142.92862 135.97372 131.89967 129.22801 127.57523 126.10547 124.90904 122.85302 120.67897 118.10106 114.94723
[15] 112.59626 110.70751 108.30899 106.38053 104.49298 103.28209 102.55376 101.90682 101.44031 100.92007 100.59363 100.40617 100.32195 100.27992
[29] 100.09197 100.09100 99.95963 99.93241 99.85075 99.67229 99.59331 99.59402 99.57548 99.51828 99.52533 99.46797 99.46126 99.44854
[43] 99.44265 99.39189 99.34614 99.21857 99.22047 99.17982 99.18918 99.19424 99.20104
> cv.r2.bwd
[1] 0.1748651 0.2411937 0.3128474 0.3625802 0.4231046 0.4571566 0.4789247 0.4921682 0.5038020 0.5131727 0.5290673 0.5455874 0.5647941 0.5877277
[15] 0.6044193 0.6175793 0.6339704 0.6468888 0.6593084 0.6671587 0.6718365 0.6759637 0.6789237 0.6822085 0.6842611 0.6854368 0.6859643 0.6862273
[29] 0.6874024 0.6874084 0.6882285 0.6883982 0.6889073 0.6900183 0.6905094 0.6905049 0.6906202 0.6909755 0.6909317 0.6912879 0.6913295 0.6914085
[43] 0.6914450 0.6917599 0.6920436 0.6928340 0.6928222 0.6930739 0.6930160 0.6929847 0.6929426
```

REGULARIZATION OR SHRINKAGE TECHNIQUES **LASSO AND RIDGE REGRESSION**



LASSO REGRESSION

```
> lasso.mod=glmnet(X[trainIndex,],y[trainIndex],alpha=1,lambda=grid,nfolds=5)
> lasso.mod

call: glmnet(x = X[trainIndex, ], y = y[trainIndex], alpha = 1, lambda = grid,
             nfolds = 5)

> cv.out.lasso

Call: cv.glmnet(x = X[trainIndex, ], y = y[trainIndex], nfolds = 5,      alpha = 1)

Measure: Mean-Squared Error
```

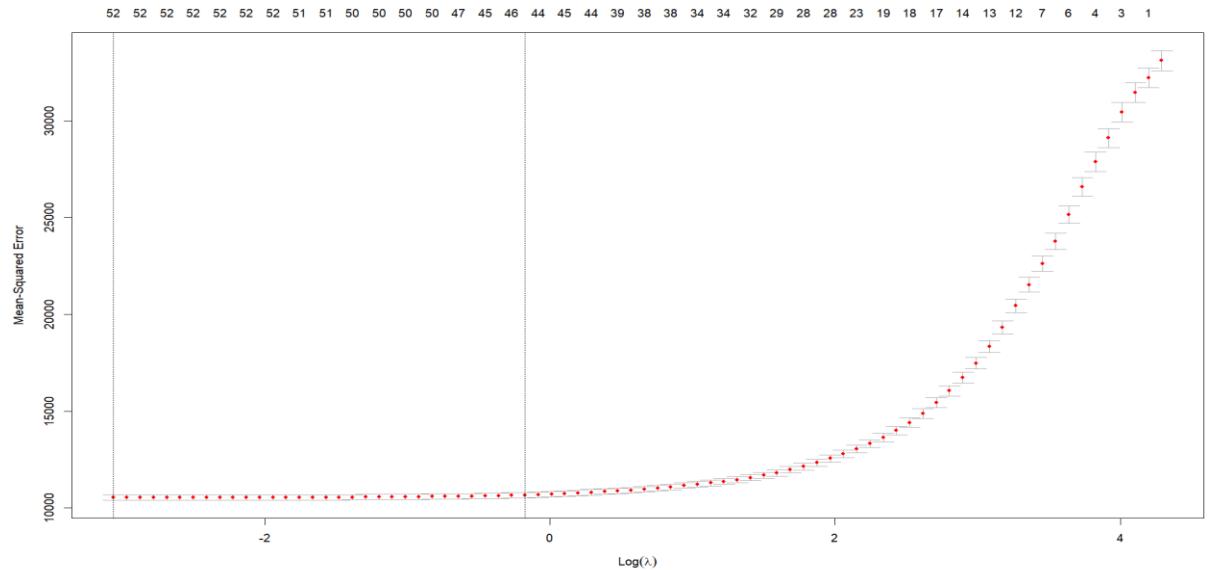
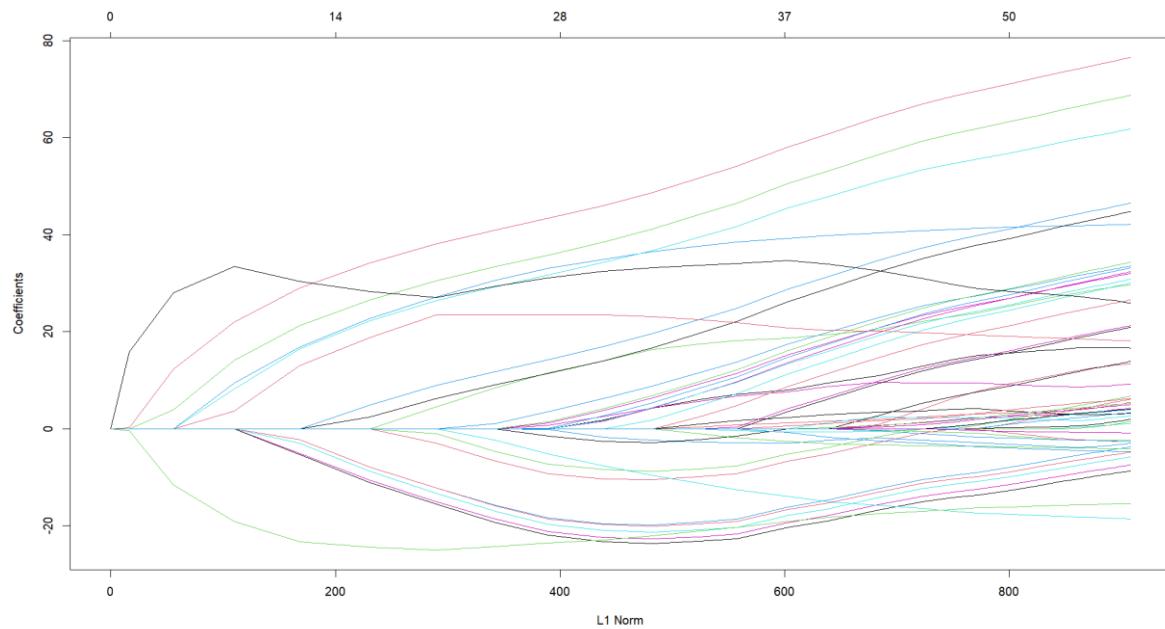
	Lambda	Index	Measure	SE	Nonzero
min	0.047	80	10561	126.5	52
1se	0.840	49	10679	139.3	45

```
> bestlam
[1] 0.04696182
```

```
> lasso.coef
(Intercept)    season2    season3    season4     yr1     mnth2     mnth3     mnth4     mnth5     mnth6     mnth7
189.4839155  16.8577497  13.6403741  28.6140927  42.6070787  0.5347265  3.4018016  0.9128913  4.8143983  0.7602159  -4.4177021
          mnth8     mnth9     mnth10    mnth11    mnth12    hr1        hr2        hr3        hr4        hr5        hr6
          1.4885979   8.3489765   4.1107508  -2.7242267  -1.8195189  -4.3940038  -6.1581060  -8.1785361  -8.7635333  -5.5658046  6.0917871
          hr7        hr8        hr9        hr10       hr11       hr12       hr13       hr14       hr15       hr16       hr17
          33.1206990  61.2071903  31.6025486  20.6256687  25.6801461  33.5438328  32.5368973  29.3447323  31.2373667  43.7272945  74.5796167
          hr18       hr19       hr20       hr21       hr22       hr23       holiday1    weekday1   weekday2   weekday3   weekday4
          68.1020591  46.3611803  30.4328204  20.5480397  13.1663565  5.4087410  -4.3636285  2.8973294  3.4436240  4.4173193  4.2494181
          weekday5    weekday6  weathersit2  weathersit3  weathersit4      temp      atemp       hum
          5.7814294   5.3450470  -4.4654339  -17.6923249  -0.7710976  23.3011725  22.0968532  -16.1323388

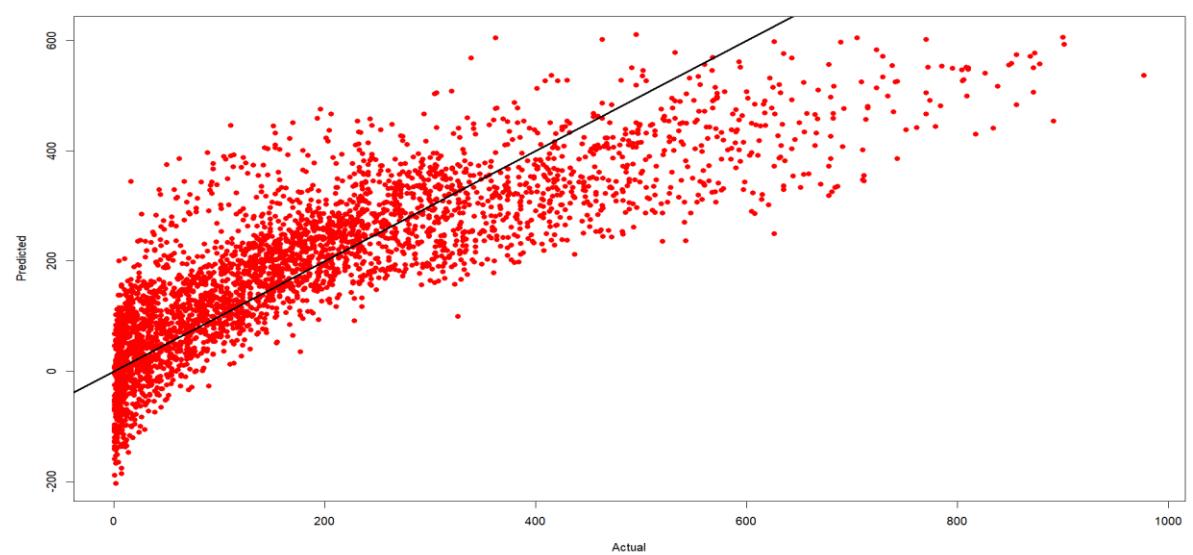
> lasso.coef[lasso.coef!=0]
(Intercept)    season2    season3    season4     yr1     mnth2     mnth3     mnth4     mnth5     mnth6     mnth7
189.4839155  16.8577497  13.6403741  28.6140927  42.6070787  0.5347265  3.4018016  0.9128913  4.8143983  0.7602159  -4.4177021
          mnth8     mnth9     mnth10    mnth11    mnth12    hr1        hr2        hr3        hr4        hr5        hr6
          1.4885979   8.3489765   4.1107508  -2.7242267  -1.8195189  -4.3940038  -6.1581060  -8.1785361  -8.7635333  -5.5658046  6.0917871
          hr7        hr8        hr9        hr10       hr11       hr12       hr13       hr14       hr15       hr16       hr17
          33.1206990  61.2071903  31.6025486  20.6256687  25.6801461  33.5438328  32.5368973  29.3447323  31.2373667  43.7272945  74.5796167
          hr18       hr19       hr20       hr21       hr22       hr23       holiday1    weekday1   weekday2   weekday3   weekday4
          68.1020591  46.3611803  30.4328204  20.5480397  13.1663565  5.4087410  -4.3636285  2.8973294  3.4436240  4.4173193  4.2494181
          weekday5    weekday6  weathersit2  weathersit3  weathersit4      temp      atemp       hum
          5.7814294   5.3450470  -4.4654339  -17.6923249  -0.7710976  23.3011725  22.0968532  -16.1323388
```

LASSO REGRESSION



Prediction with test data

```
> # Display the results
> cat("R-squared: ", r_squared_lr, "\n")
R-squared: 0.6929642
> cat("Mean Square Error (MSE): ", mse_lr, "\n")
Mean Square Error (MSE): 9840.154
> cat("Root Mean Square Error (RMSE): ", rmse_lr, "\n")
Root Mean Square Error (RMSE): 99.19755
```



RIDGE REGRESSION

```
> ridge.mod=glmnet(X[trainIndex,],y[trainIndex],alpha=0,lambda=grid,nfold = 5)
> ridge.mod

Call: glmnet(x = X[trainIndex, ], y = y[trainIndex], alpha = 0, lambda = grid,
             nfold = 5)

> cv.out

Call: cv.glmnet(x = X[trainIndex, ], y = y[trainIndex], nfolds = 5,      alpha = 0)

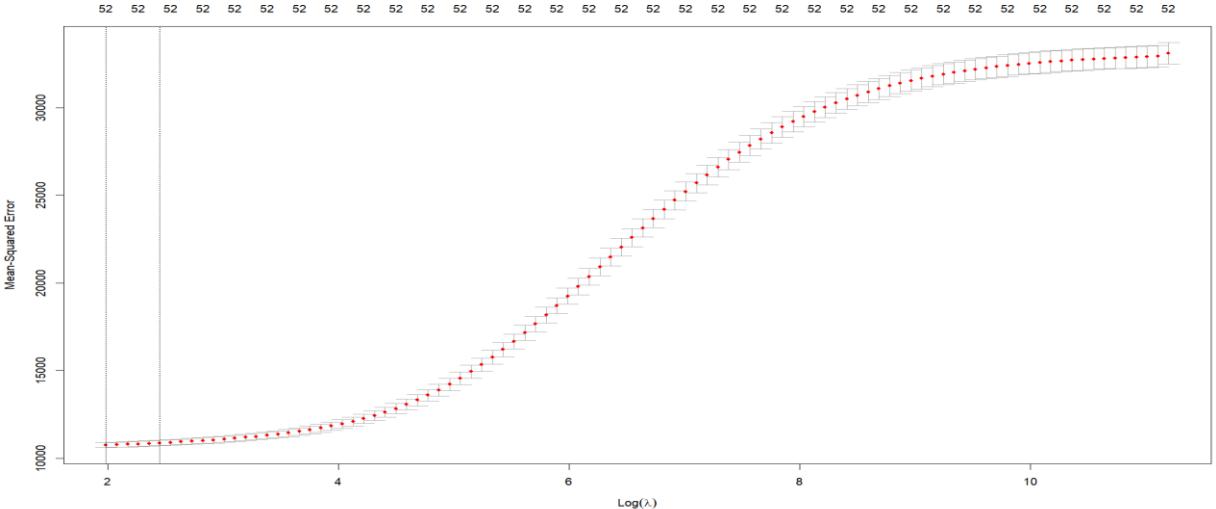
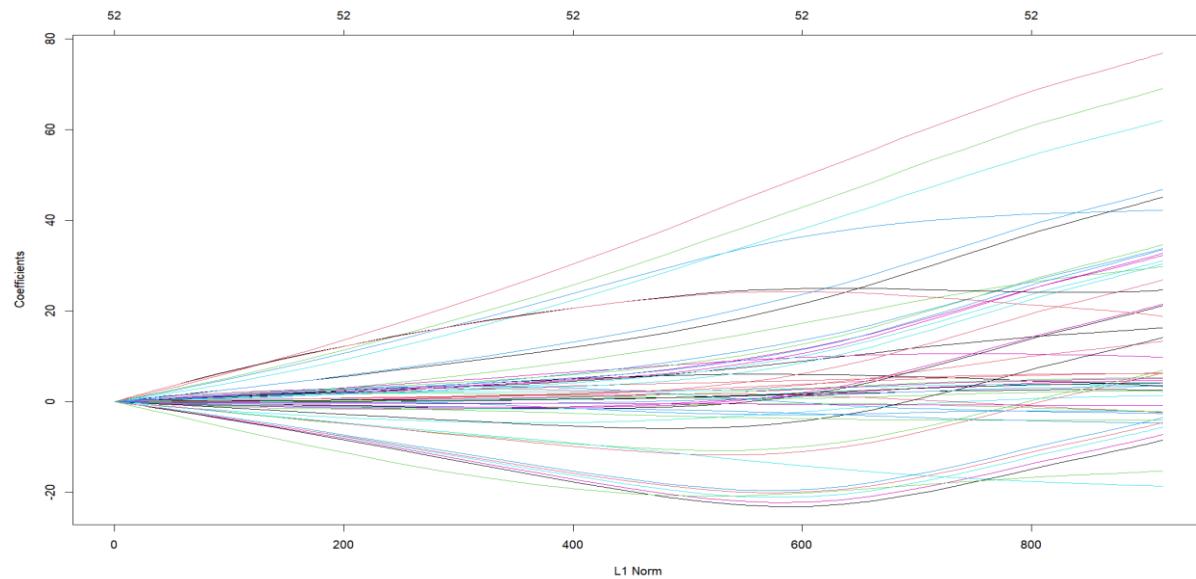
Measure: Mean-Squared Error

      Lambda Index Measure    SE Nonzero
min   7.306   100 10752 146.1      52
1se  11.633    95 10871 157.6      52
```

> bestlam
[1] 7.305738

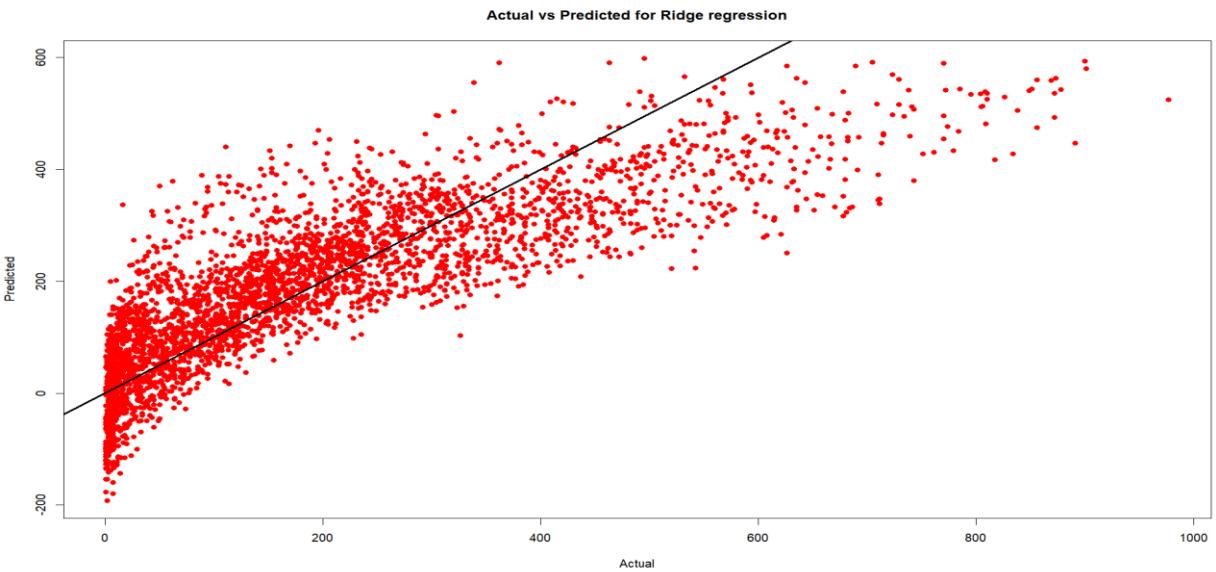
```
> ridge.coef
(Intercept)  season2  season3  season4  yr1  mnth2  mnth3  mnth4  mnth5  mnth6  mnth7  mnth8  mnth9
189.4839155 12.5550578 7.7016991 22.8651171 40.6945589 -0.4186437 2.9264633 1.8598086 5.6008711 1.4196220 -3.1470141 2.9491512 9.9927804
mnth10  mnth11  mnth12  hr1  hr2  hr3  hr4  hr5  hr6  hr7  hr8  hr9  hr10
6.2129840 -0.1525871 -0.2738258 -14.8754593 -16.4521741 -18.2137968 -18.7318176 -15.7770654 -4.5358967 21.3845958 48.3075871 19.6213550 8.8735587
hr11  hr12  hr13  hr14  hr15  hr16  hr17  hr18  hr19  hr20  hr21  hr22  hr23
13.5712367 20.9766971 19.8998662 16.7523456 18.5526797 30.5976738 60.3741305 54.2571669 33.4940456 18.2856256 8.9042161 1.8890558 -5.5148611
holiday1  weekday1  weekday2  weekday3  weekday4  weekday5  weekday6  weathersit2  weathersit3  weathersit4  temp  atemp  hum
-4.4042221 2.3121205 2.8247142 3.7550335 3.5338167 5.0127536 4.6124441 -3.4852771 -15.9546090 -0.6671070 24.2334388 23.9738209 -18.6403522
> ridge.coef[ridge.coef!=0]
(Intercept)  season2  season3  season4  yr1  mnth2  mnth3  mnth4  mnth5  mnth6  mnth7  mnth8  mnth9
189.4839155 12.5550578 7.7016991 22.8651171 40.6945589 -0.4186437 2.9264633 1.8598086 5.6008711 1.4196220 -3.1470141 2.9491512 9.9927804
mnth10  mnth11  mnth12  hr1  hr2  hr3  hr4  hr5  hr6  hr7  hr8  hr9  hr10
6.2129840 -0.1525871 -0.2738258 -14.8754593 -16.4521741 -18.2137968 -18.7318176 -15.7770654 -4.5358967 21.3845958 48.3075871 19.6213550 8.8735587
hr11  hr12  hr13  hr14  hr15  hr16  hr17  hr18  hr19  hr20  hr21  hr22  hr23
13.5712367 20.9766971 19.8998662 16.7523456 18.5526797 30.5976738 60.3741305 54.2571669 33.4940456 18.2856256 8.9042161 1.8890558 -5.5148611
holiday1  weekday1  weekday2  weekday3  weekday4  weekday5  weekday6  weathersit2  weathersit3  weathersit4  temp  atemp  hum
-4.4042221 2.3121205 2.8247142 3.7550335 3.5338167 5.0127536 4.6124441 -3.4852771 -15.9546090 -0.6671070 24.2334388 23.9738209 -18.6403522
```

RIDGE REGRESSION



Prediction with test data

```
> cat("R-squared: ", r_squared_rr, "\n")
R-squared: 0.6877307
> cat("Mean Square Error (MsE): ", mse_rr, "\n")
Mean Square Error (MsE): 10007.88
> cat("Root Mean Square Error (RMSE): ", rmse_rr, "\n")
Root Mean Square Error (RMSE): 100.0394
```

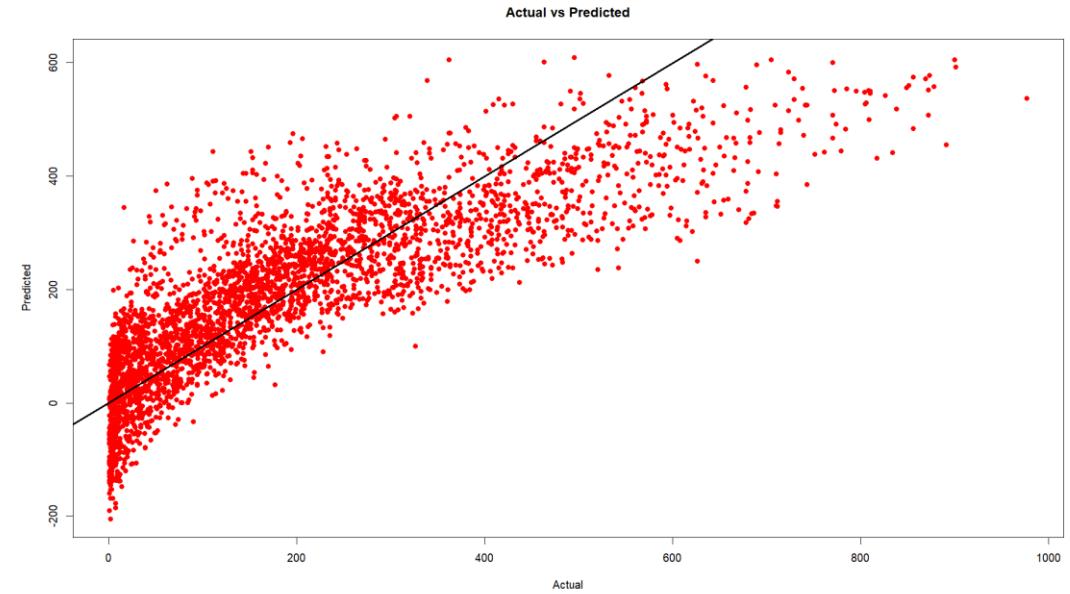
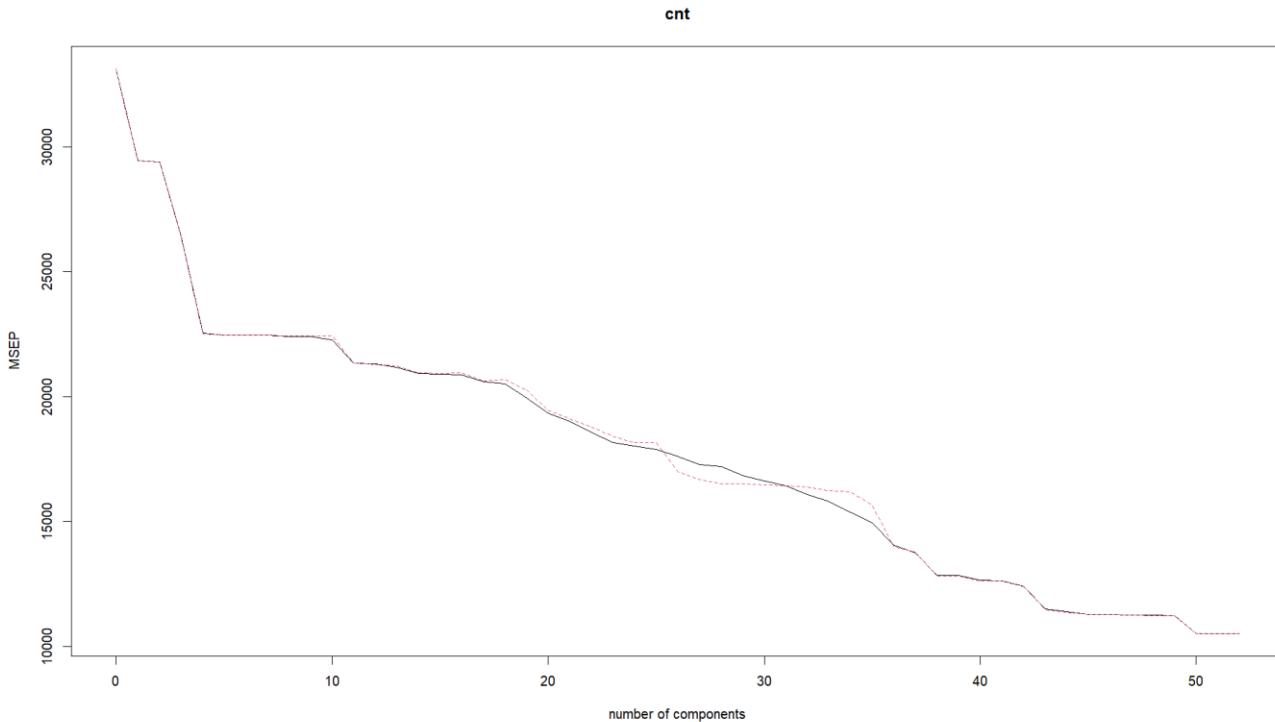




DIMENSION REDUCTION METHOD PRINCIPLE COMPONENT REGRESSION

PRINCIPLE COMPONENT REGRESSION

```
> pcr.fit
Principal component regression, fitted with the singular value decomposition algorithm.
Cross-validated using 10 random segments.
Call:
pcr(formula = cnt ~ ., data = dff, subset = trainIndex, scale = TRUE,      validation = "CV")
> summary(pcr.fit)
Data: X dimension: 13902 52
Y dimension: 13902 1
Fit method: svdpc
Number of components considered: 52
> cat("R-squared: ", r_squared_pca, "\n")
R-squared: 0.692995
> cat("Mean Square Error (MsE): ", mse_pca, "\n")
Mean Square Error (MsE): 9839.166
> cat("Root Mean Square Error (RMSE): ", rmse_pca, "\n")
Root Mean Square Error (RMSE): 99.19257
```





RANDOM FOREST



RANDOM FOREST

```
> rf.model
```

Call:

```
randomForest(formula = cnt ~ ., data = dff, importance = TRUE, subset = trainIndex)
```

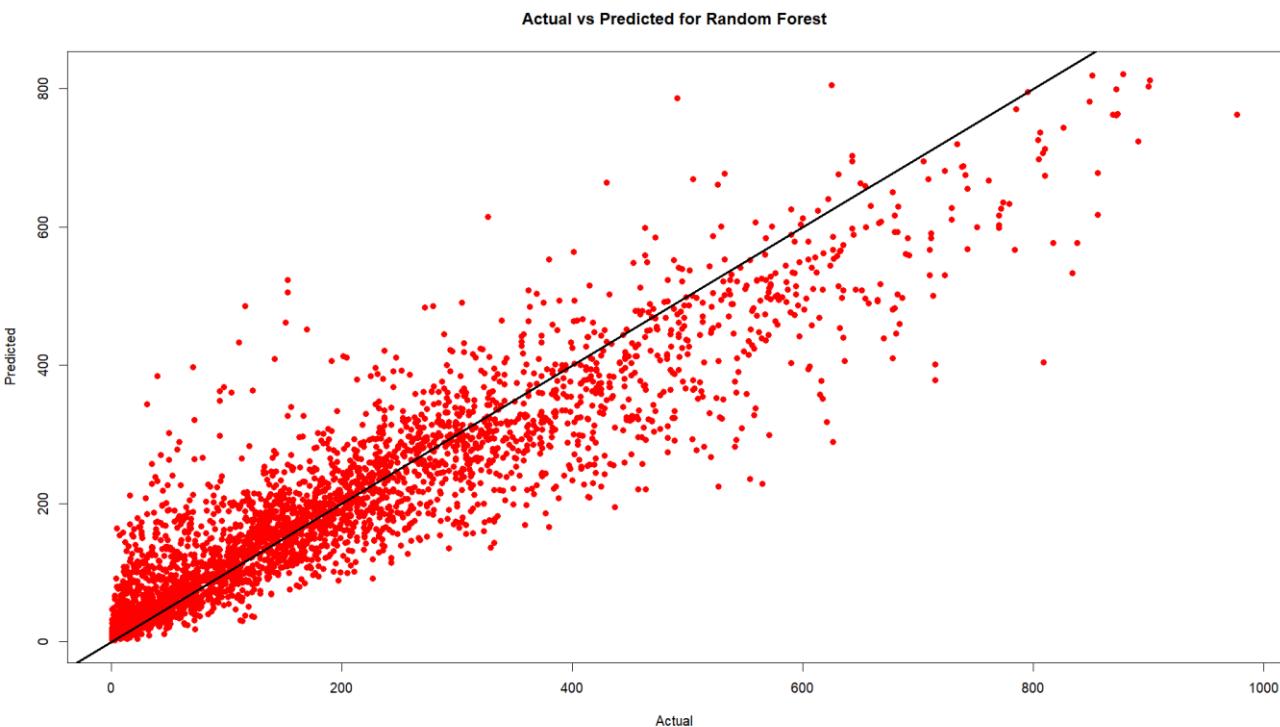
Type of random forest: regression

Number of trees: 500

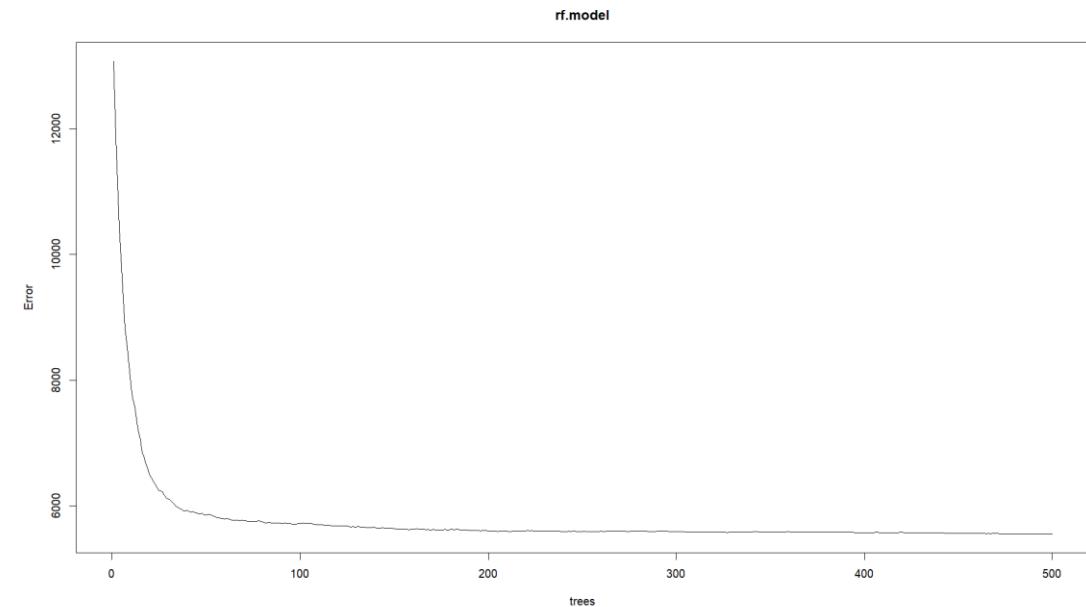
No. of variables tried at each split: 17

Mean of squared residuals: 5555.428

% Var explained: 83.22



Trees Vs Error



```
> cat("R-squared: ", r_squared_rf, "\n")  
R-squared: 0.8443822  
> cat("Mean Square Error (MSE): ", mse_rf, "\n")  
Mean Square Error (MsE): 4987.375  
> cat("Root Mean Square Error (RMSE): ", rmse_rf, "\n")  
Root Mean Square Error (RMSE): 70.62135
```



BOOSTING TECHNIQUE GRADIENT BOOSTING MACHINE

GBM

```
> boost.model=gbm(cnt~., data=dff[trainIndex,], distribution="gaussian", n.trees=5000, cv.folds = 5, interaction.depth=6)
> boost.model
```

```
gbm(formula = cnt ~ ., distribution = "gaussian", data = dff[trainIndex,
  ], n.trees = 5000, interaction.depth = 6, cv.folds = 5)
```

A gradient boosted model with gaussian loss function.

5000 iterations were performed.

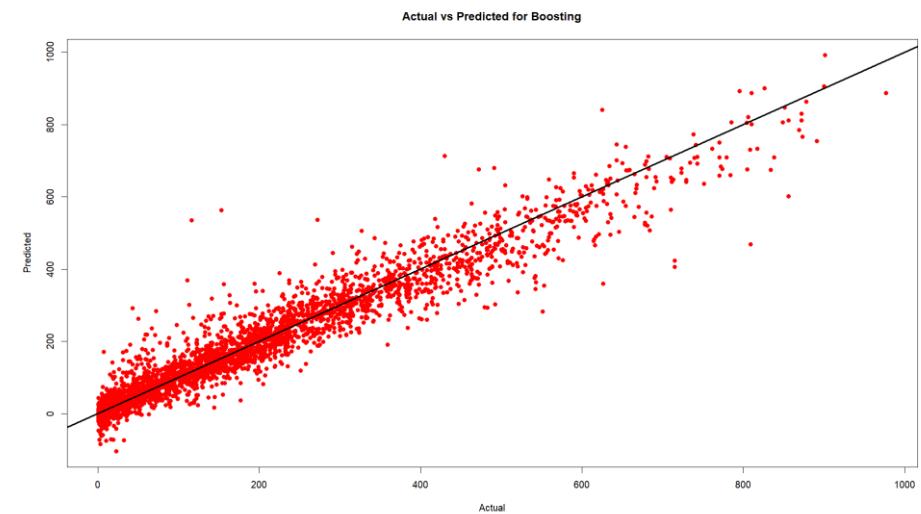
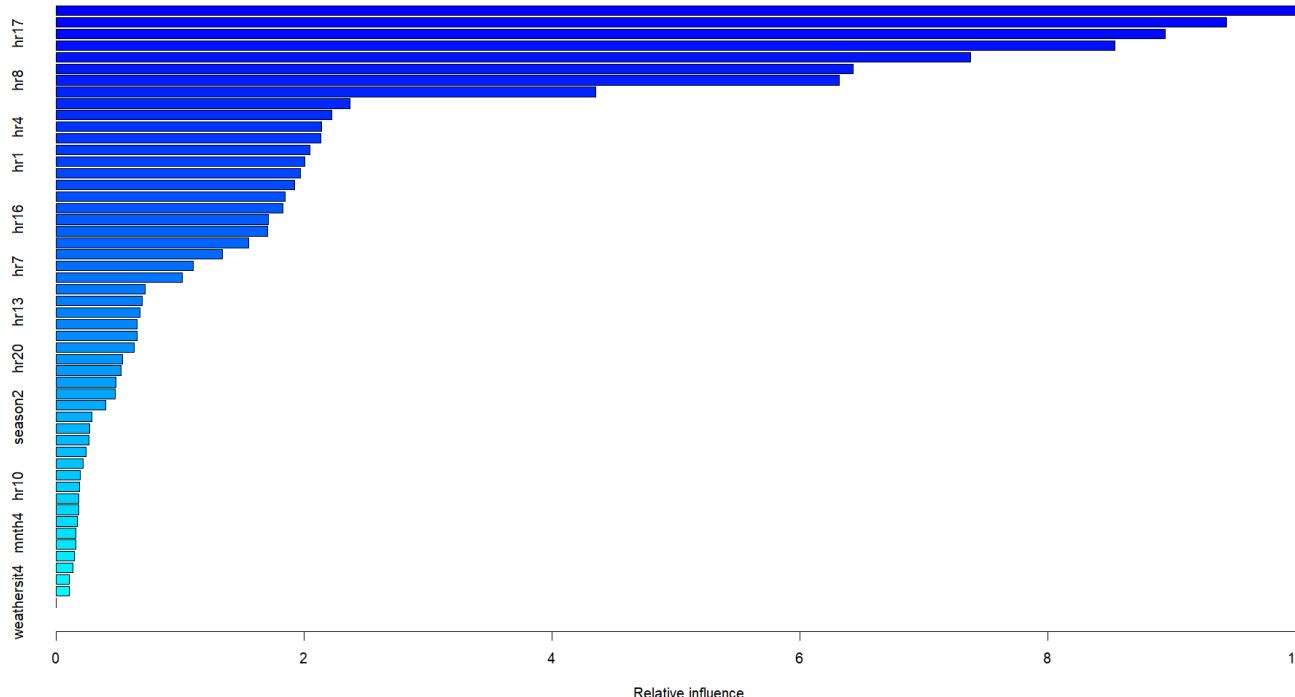
The best cross-validation iteration was 4806.

There were 52 predictors of which 51 had non-zero influence.

```
> cat("R-squared: ", r_squared_bm, "\n")
R-squared: 0.9289999
```

```
> cat("Mean Square Error (MsE): ", mse_bm, "\n")
Mean Square Error (MsE): 2275.473
```

```
> cat("Root Mean Square Error (RMSE): ", rmse_bm, "\n")
Root Mean Square Error (RMSE): 47.70192
```

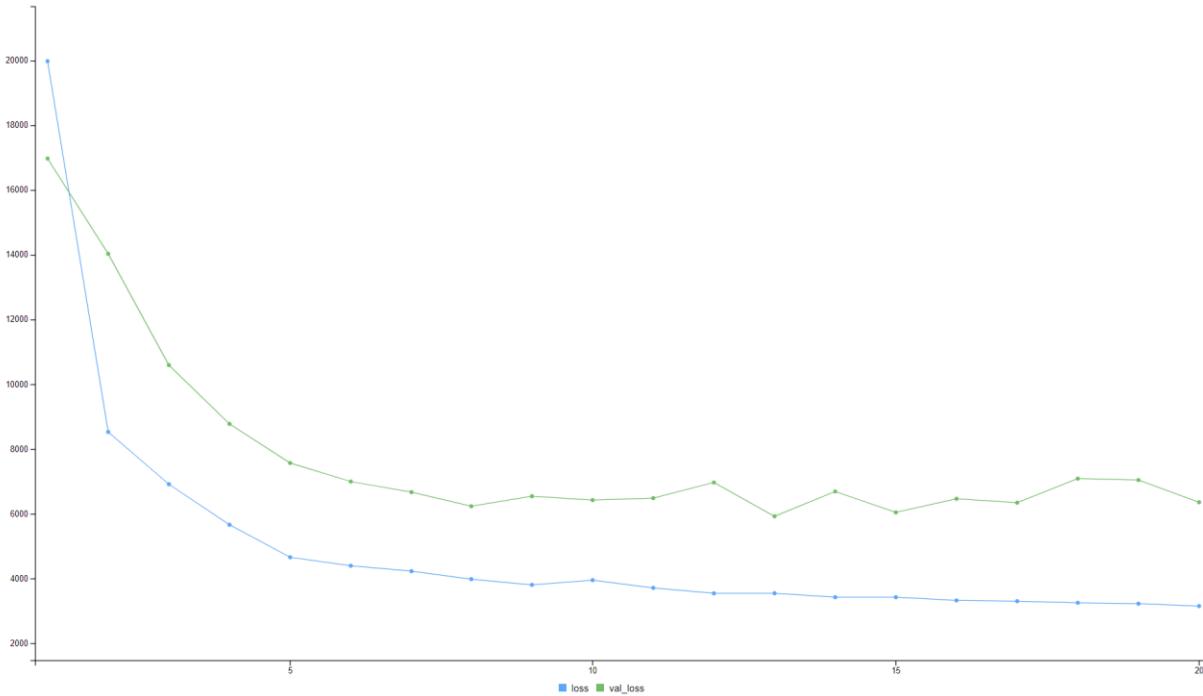




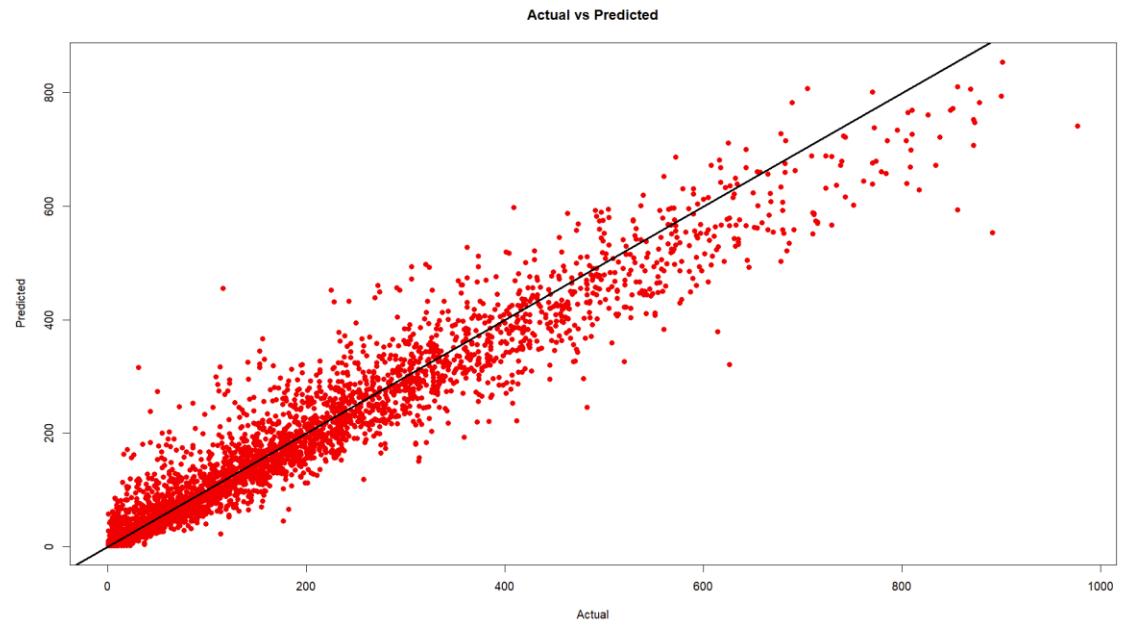
NEURAL NETWORK

NEURAL NETWORK

- A neural network is a computational model inspired by the brain's biological networks, using interconnected nodes or neurons to process and learn from data. It excels in pattern recognition and problem-solving in various AI and machine learning applications.



```
> model
Model: "sequential_17"
-----  
Layer (type)          Output Shape       Param #  
=====  
dense_69 (Dense)      (None, 128)        6784  
dropout_43 (Dropout)   (None, 128)        0  
dense_68 (Dense)      (None, 64)         8256  
dropout_42 (Dropout)   (None, 64)         0  
dense_67 (Dense)      (None, 32)         2080  
dropout_41 (Dropout)   (None, 32)         0  
dense_66 (Dense)      (None, 1)          33  
=====  
Total params: 17153 (67.00 KB)  
Trainable params: 17153 (67.00 KB)  
Non-trainable params: 0 (0.00 Byte)
```

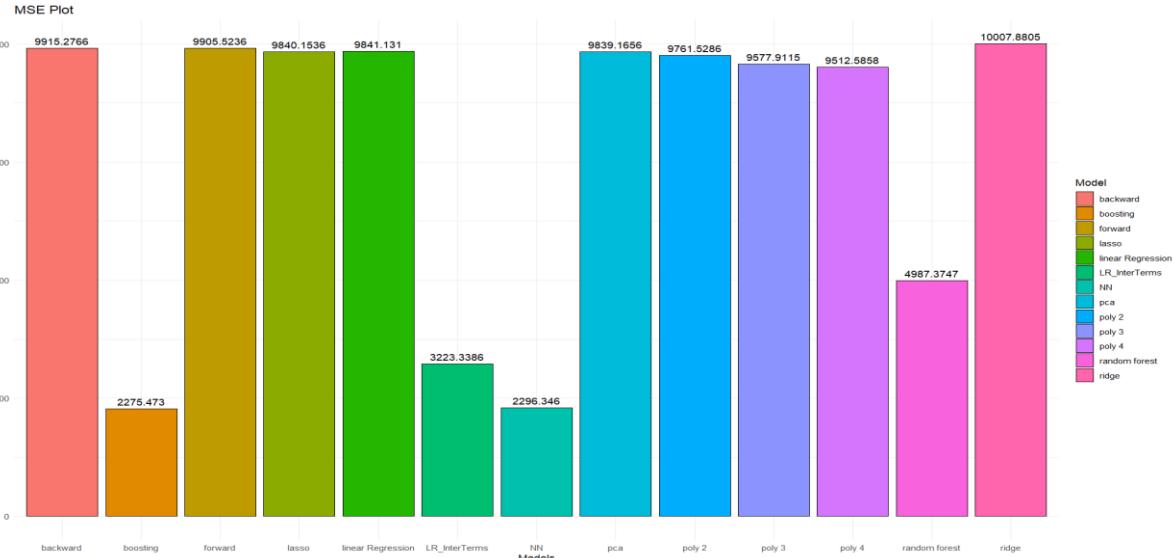


```
> npred[90:120]
[1] 53.524258 100.834198 95.815201 54.177750 16.889029 14.479589 6.393068 59.633827 47.562016
[10] 34.576405 34.139233 159.828445 122.070030 4.273767 28.845617 53.216862 10.493546 88.245804
[19] 63.974056 37.307816 73.467369 37.573139 11.945144 8.971384 38.799953 68.332489 155.409195
[28] 52.531227 21.633905 43.107937 97.458145
> y_test[90:120]
[1] 79 104 118 67 17 8 2 47 22 28 35 125 133 2 52 59 6 143 98 81 91 56 21 2 63 65
[27] 161 53 11 40 70
```

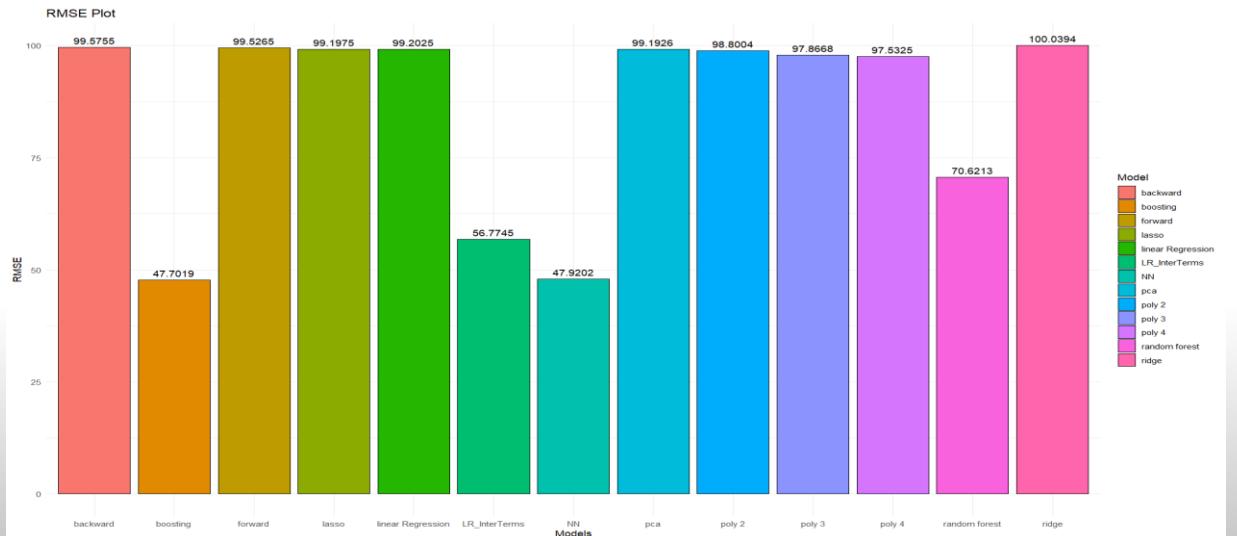
```
> cat("R-squared: ", r_squared_nn, "\n")
R-squared: 0.9231235
> cat("Mean Square Error (MsE): ", mse_nn, "\n")
Mean Square Error (MsE): 2463.806
> cat("Root Mean Square Error (RMSE): ", rmse_nn, "\n")
Root Mean Square Error (RMSE): 49.63674
>
```

NEURAL NETWORK

RESULTS

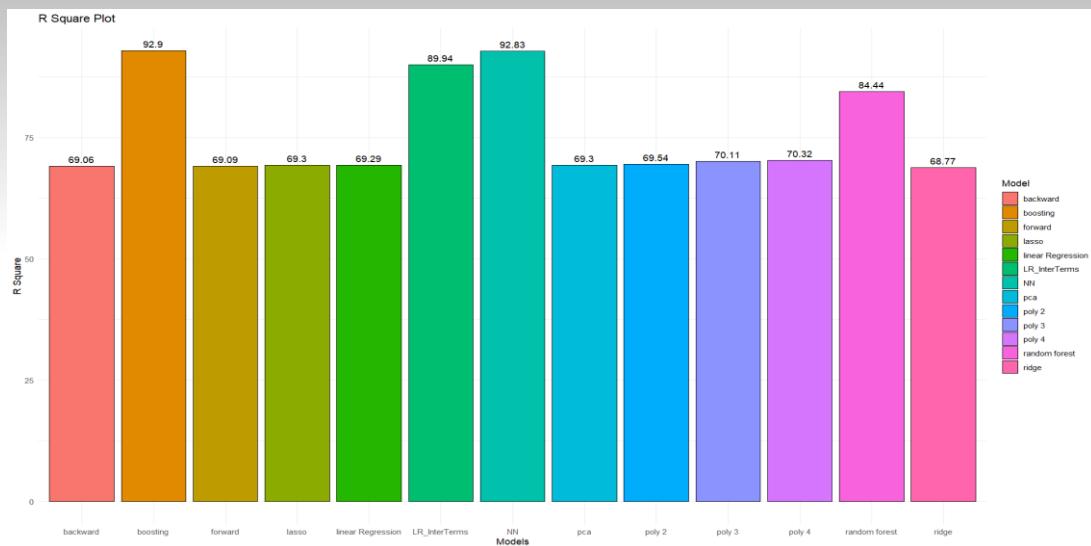


Mean Square Error (MSE)



Root Mean Square Error (RMSE)

R Square (R^2)



MODEL EVALUATION

Model	MSE	RMSE	R ²
Linear Regression	9841.131	99.20247	69.29337
Linear Regression – Poly 2	9761.529	98.80045	69.54175
Linear Regression – Poly 3	9577.911	97.8668	70.11467
Linear Regression – Poly 4	9512.586	97.53249	70.31851
Linear Regression – Interaction term	3223.339	56.77445	89.94243
Forward Sub Selection	9905.524	99.5265	69.09245
Backward Sub selection	9915.277	99.57548	69.06202
Lasso Regression	11595.847	107.68402	63.81824
Ridge Regression	9980.374	99.90182	68.8589
Random Forest	4987.375	70.62135	84.43822
Gradient Boosting Machine	2275.473	47.70192	92.89999
Principle Component Analysis	9839.166	99.19257	69.2995
Neural Network	2362.696	48.60757	92.62784

CONCLUSION

- The Gradient Boosting Machine (GBM) and Neural Network models outperformed others in predicting bike-sharing demand, with the GBM showing the lowest error rates and a high R^2 value, highlighting its precise predictions.
- The elapsed time for Neural Network is good when compared to GBM, RF, Linear Regression with interaction terms.
- While traditional models like Linear Regression lagged in performance, introducing interaction terms significantly improved results, underscoring the importance of feature interplay in predictive accuracy.
- The Random Forest algorithm also demonstrated solid predictive capabilities, suggesting it as a reliable alternative. Overall, the GBM and Neural Network are the preferred models for their superior prediction quality in bike-sharing demand forecasting.

FUTURE SCOPE

- In the future, our project aims to significantly expand its scope and impact. We plan to integrate real-time data to enhance prediction accuracy, and extend the model to encompass a wider range of urban transportation modes, providing a more comprehensive view of city mobility. Delving deeper into advanced machine learning techniques, we'll strive to refine our predictive capabilities, ensuring the model adapts to diverse urban environments globally.
- Additionally, by analyzing user behavior more closely, we aim to tailor the bike-sharing experience, increasing its efficiency and appeal. These steps are not just about advancing the project, but also about contributing to larger goals like reducing urban congestion and promoting sustainable transportation in cities around the world.